


Getting up to speed @AIRLab



Things are moving fast, so this presentation might be outdated and/or incomplete.
Contribution and improvements are welcome!

Copyrights © 2023 Nico Catalano
Contacts: nico.catalano@polimi.it

Disclaimer

Most of the concepts in this presentation are just suggestions for a great start

-> It can't contain everything you will ever need

Every thesis is different, **you will learn to learn by yourself!**

Who told you what to do during your thesis ?



No one



Tesista's check list



Make sure to be able to access the **building 7**

<https://airlab.deib.polimi.it/life-in-airlab/service-pages/becoming-an-airlab-user/>



Ask your co-advisor for the **VPN** credentials if not already provided



Ask your co-advisor for the servers' **credentials** if not already provided



Ask your co-advisor the **AIRLab Discord** server joining link if not already provided

How to approach the thesis' challenges

- The thesis is **YOUR PROJECT**, so it should reflect **YOUR INTEREST** and ambition
- The co-advisor is there to help you getting there
- You should enjoy the ride
- There is not a fixed time frame
- **Do not hesitate to talk to your co-advisor!**



Things to know before start writing

- Notes on how to write the thesis:

<https://airlab.deib.polimi.it/wp-content/uploads/2021/03/How-to-write-a-thesis.pdf>

- Rules:

<https://www.ingindinf.polimi.it/en/1/teaching/lectures-and-exams/degree-examinations>

- POLIMI's LATEX (Overleaf) template:


<https://www.overleaf.com/gallery/tagged/polimi>

- POLIMI's MS PowerPoint template:

https://polimi365-my.sharepoint.com/:f/g/personal/10306579_polimi_it/Eu7UoRKAz-pHmHKuXXwdFd8Bn4JVyBfeYbMOm3uIGg4pkg?e=zLNdnj

General Suggestions

- Take notes during meetings
- Schedule regular update meeting with your co-advisor
- Always prepare what to show and talk about for scheduled meetings
- Keep track of your readings (eg using Zotero)
- **Do not use Latex for presentations, only for your thesis!**



The background of the slide is decorated with various abstract geometric shapes. In the top left, there is a blue circle and a green triangle. Below the triangle is a green square. To the left of the square are two yellow vertical bars. In the center, there is a large orange semi-circle and a blue circle. At the bottom, there is a large orange circle surrounded by several yellow curved lines.

AIRLab Network

AIRLab's servers




westworld – GPU Server

8 x 11GB Nvidia 1080 TI
40 core
252 GB RAM



elysium – GPU Server

4 x 24GB Nvidia RTX6000
1 x 12GB Nvidia Titan X
1 x 11GB Nvidia 1080TI
48 core
252 GB RAM



magrathea – GPU Server


8 x 24GB Nvidia RTX6000
112 core
252 GB RAM




multiverse – NAS

Notice! To obtain the servers IPs, please ask your co-supervisor.


AIRLab's servers



westworld – GPU Server



elysium – GPU Server



magrathea – GPU Server




multiverse – NAS

Computation only

Storage

AIRLab's servers



westworld – GPU Server



elysium – GPU Server



magrathea – GPU Server



multiverse – NAS

Students

PhDs and Postdocs

Computation only

Storage

Your account on the servers

- Same user/password on westworld, elysium (and magrathea)
- Home directory in each server:

 /surname

 /dataset

 /shared -> link to multiverse


 /private -> link to multiverse

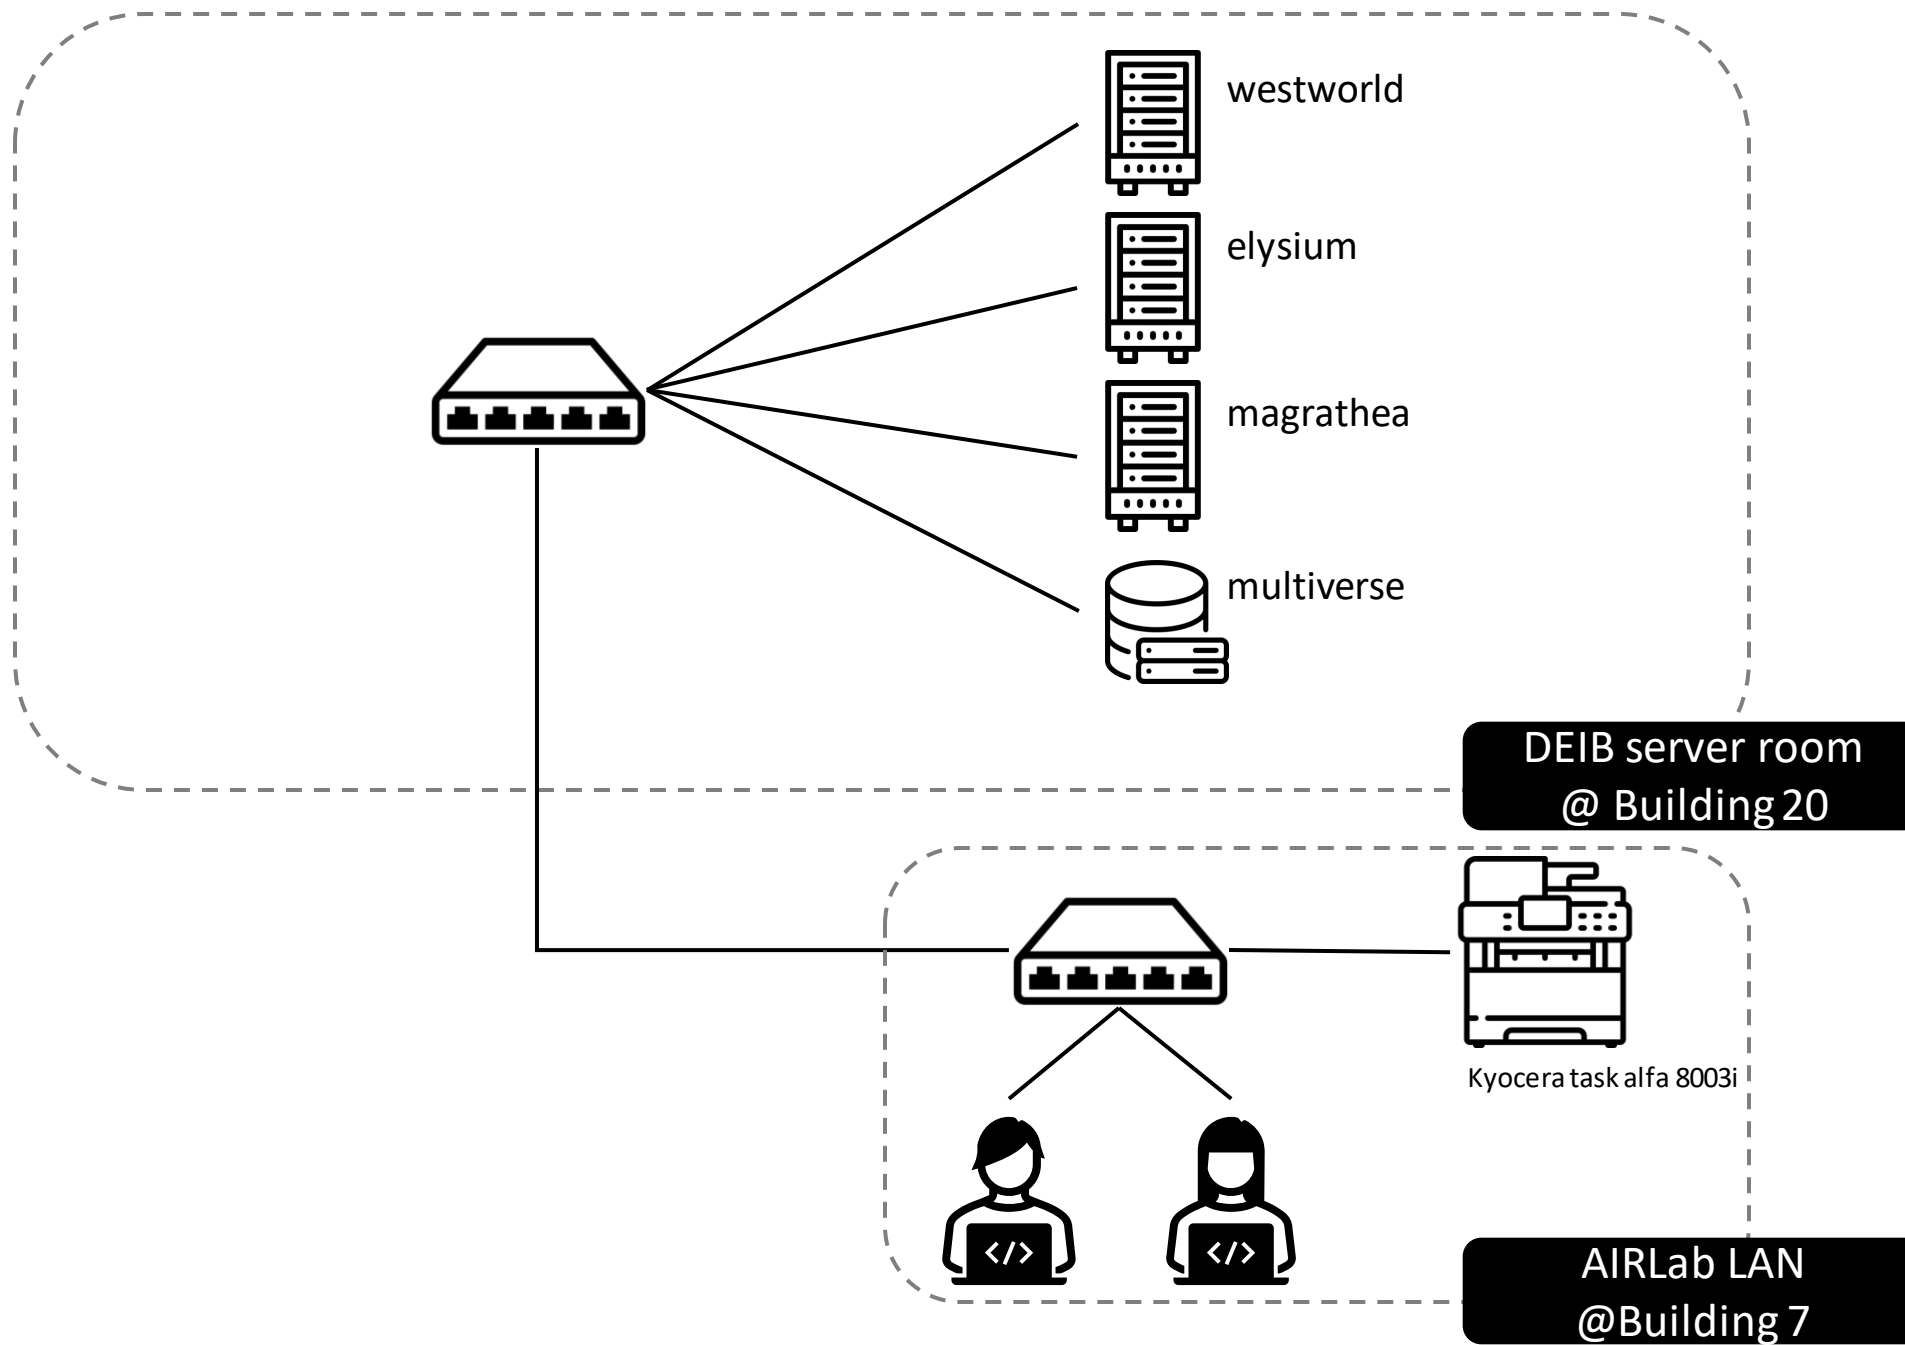
 /storage -> link to multiverse

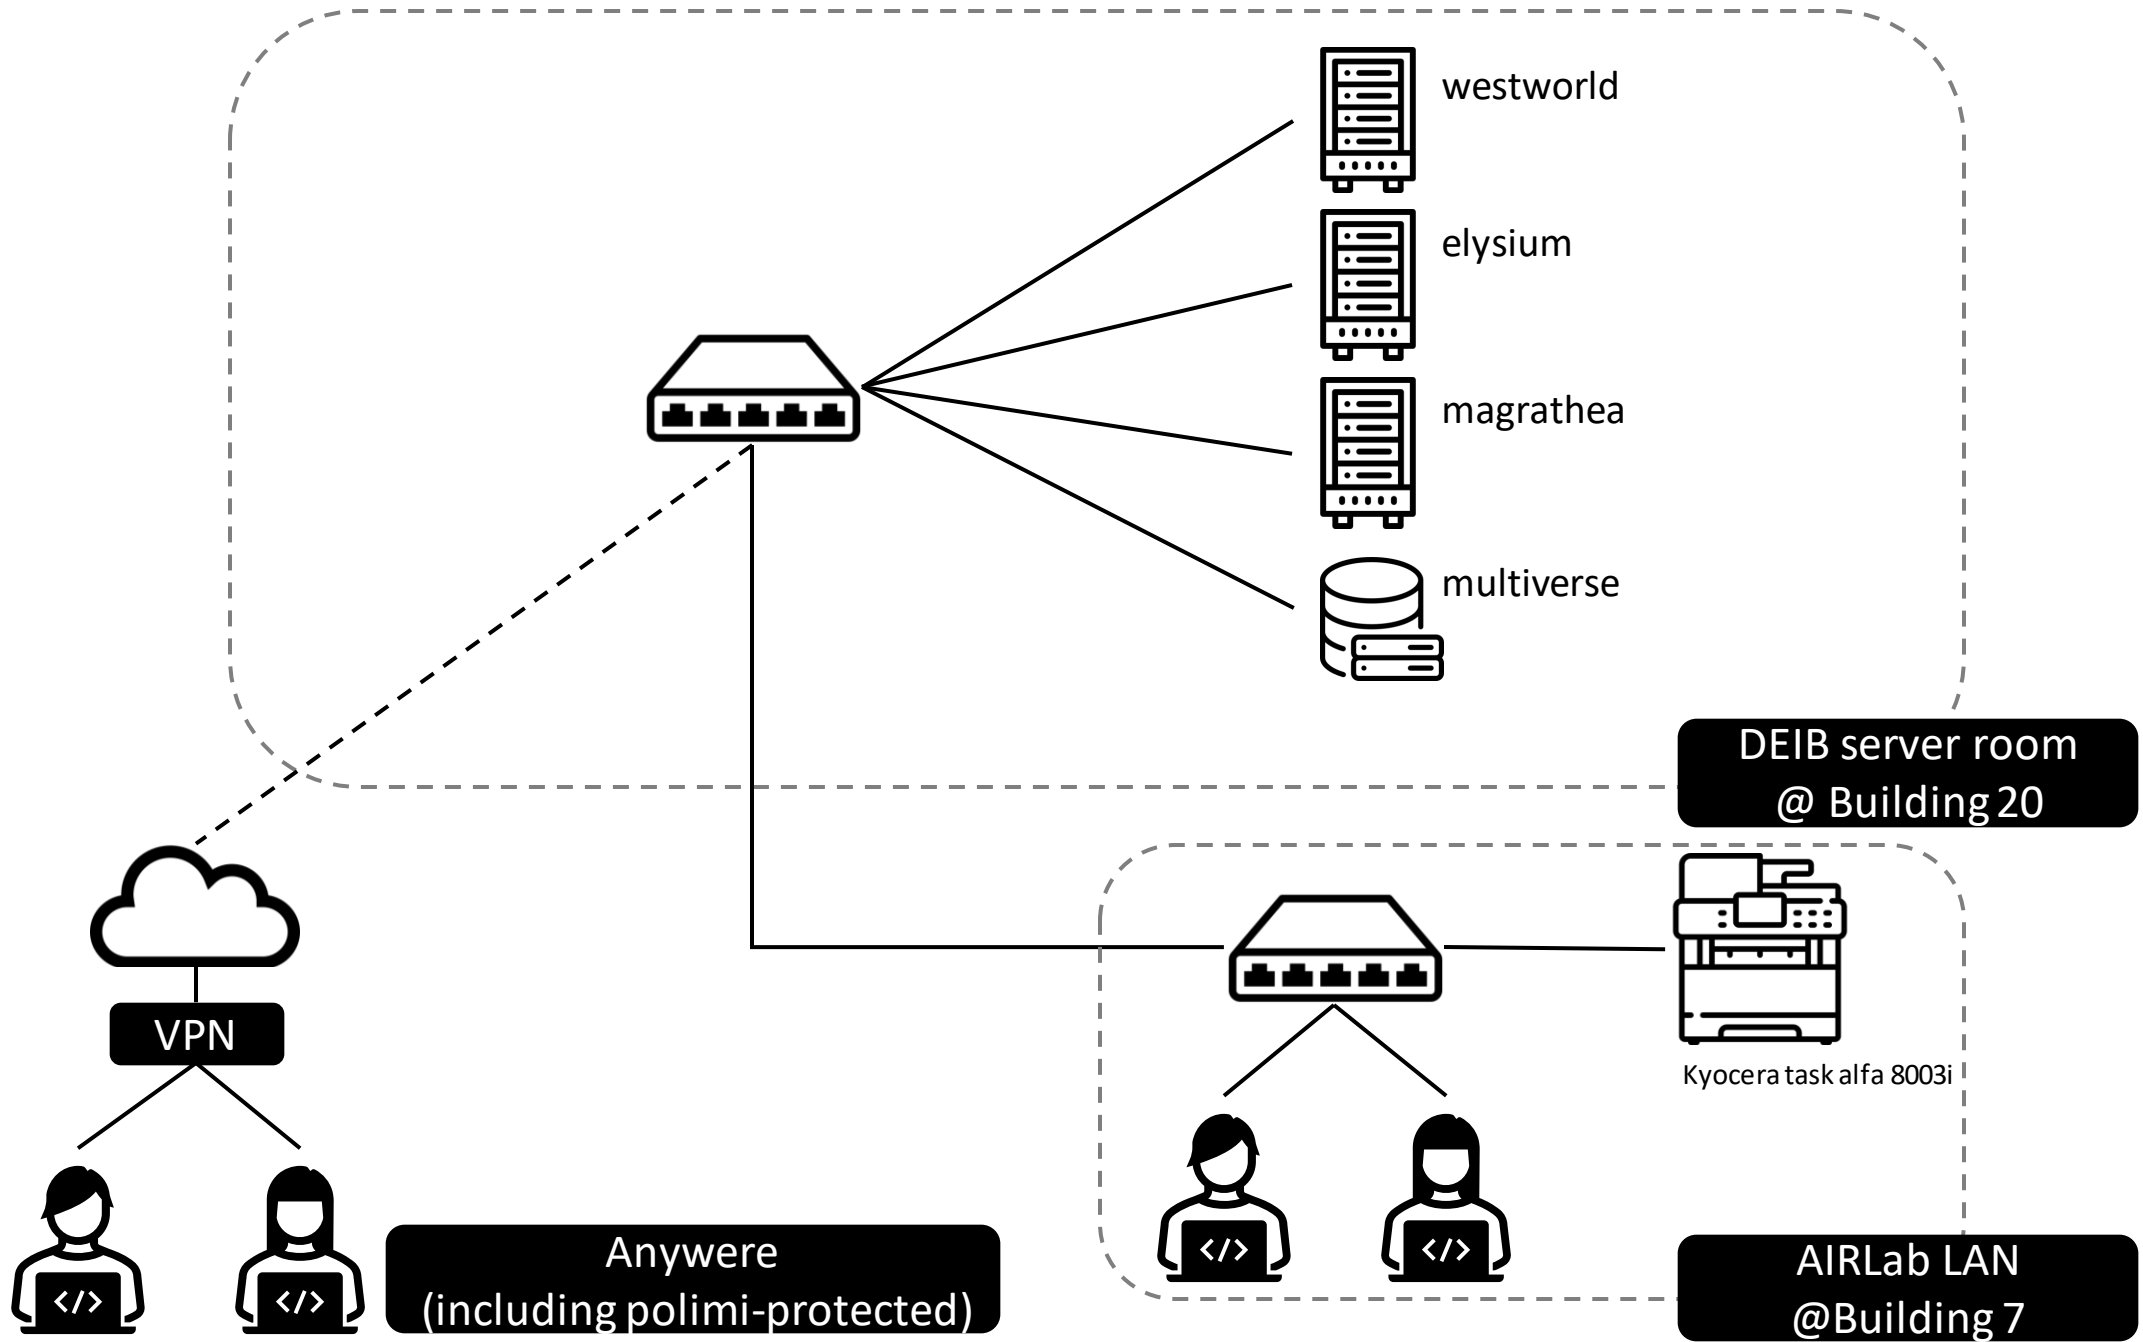
Your home folder is supposed to be empty besides the link you already find. Please put your code under /surname/storage and the datasets in /surname/dataset/private

What is on multiverse is visible from all servers

AIRLab's network







The left side of the image features several abstract geometric shapes. At the top left is a blue circle. Below it are two short, vertical yellow dashes. To the right of these is a large orange semi-circle. Further right is a green triangle pointing downwards. Below the semi-circle is a blue circle. At the bottom left is a green square. To its right are three short, curved yellow dashes. At the bottom center is a large orange circle.

Use and share
resources

Run experiments

- Share resources (GPU, CPU, ..)
- Virtualization
- Avoid dependencies conflicts
- Easier reproducibility



If you are not familiar with docker is advisable to check out some introductory tutorial on line

Book resources

Before using any GPU or CPU check that is not already used and book it for you

- Westworld resources:
https://docs.google.com/spreadsheets/d/1n6HDbSX0Pe0zcRA0iBYDStrud7d_yRhBSrpFld4WNxs/edit#gid=1311829678
- Elysium resources:
https://docs.google.com/spreadsheets/d/1wvzz0ZgPypepkZJtZahAmIJxV-iOIhsR7_s5Mht4l1A/edit#gid=0
- Magrathea resources:
<https://docs.google.com/spreadsheets/d/1zpmh4hSZp8u45HZ4p2aKrMqlBsSQHbYEthdlzRkKj2o/edit#gid=0>

Tutorial

- Get ssh certificate for servers
- Configure ssh autocomplete
- Configure network drive on ubuntu
- Configure ssh certificate for github on server
- Tmux
- Nvidia-smi
- Introduction to docker
- “Tutorial” pytorch + jupyter + tensorboard

Tutorial

- The following tutorial assumes you are using Ubuntu as operative system and VS Code as editor
- It mainly cover how I have configurated my machine -> there might be better ways, please tell me about it
- You can use whatever OS/Editor you prefer



Configure VPN

If you can't connect to the AIRLab ethernet network, you can always connect to any of the servers using the VPN.

Instructions for OSX and Windows on POLIMI web site at

<https://www.ict.help.polimi.it/network-vpn-global/>

VPN installation on Ubuntu

```
sudo apt install network-manager-openconnect network-manager-openconnect-gnome -y
```

Extract certificates (use the unique password you created the certificate with):
(with newer openssl versions, from 3.0, append the “-legacy” option to the following commands)

```
openssl pkcs12 -in YOUR_CERTIFICATE_FILE.p12 -out usercert.pem -clcerts -nokeys  
openssl pkcs12 -in YOUR_CERTIFICATE_FILE.p12 -out userkey.pem -nocerts -nodes  
openssl pkcs12 -in YOUR_CERTIFICATE_FILE.p12 -cacerts -nokeys -chain -out userca.pem
```

Work from home.



Configure in GNOME

Under settings:

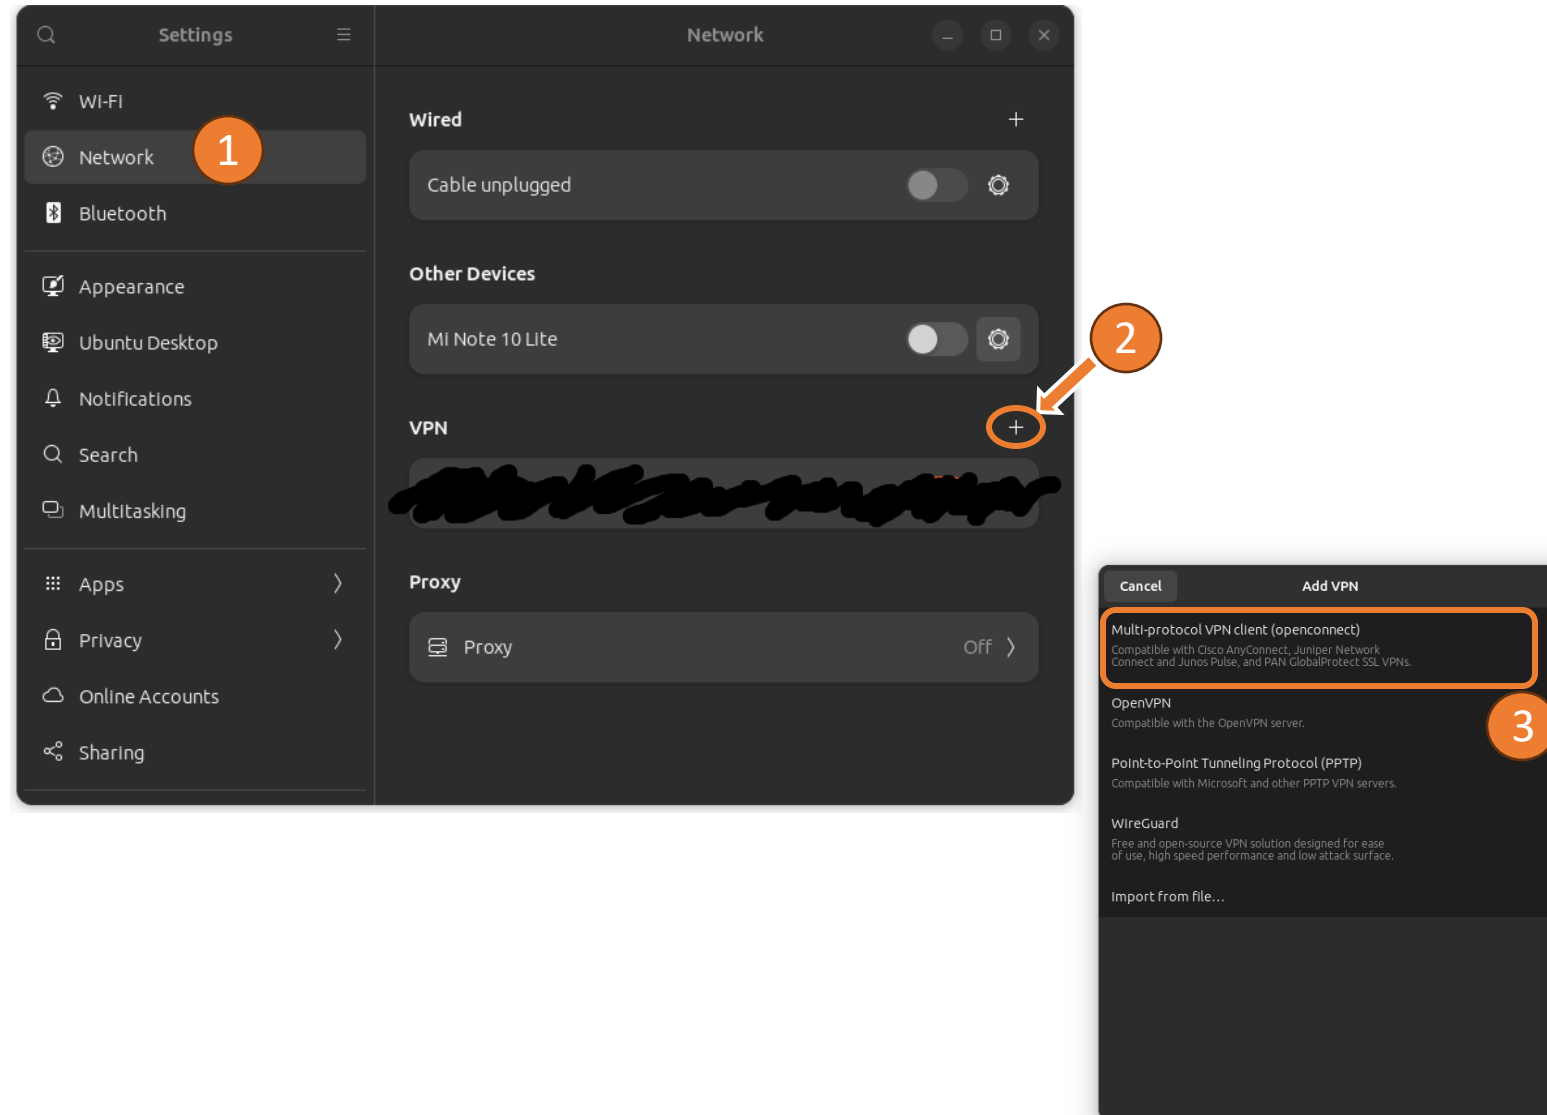
1. Network
2. Add VPN
3. Multi-protocol VPN client (openconnect)

Fill with:

- VPN Protocol: Palo Alto Networks GlobalProtect
- Gateway: gp-deib.vpn.polimi.it
- CA Certificate: userca.pem
- User Certificate: usercert.pem
- Private Key: userkey.pem

Launch the vpn from the menu filling in your polimi login credentials

([name.surname@\(mail.\)polimi.it](mailto:name.surname@(mail.)polimi.it) and password)



Connect to the servers

... Or how I configured my ssh client

Names instead of IP

- Instead of typing each time the IP of the server you want to connect to -> **use DNS entry**
- Add the following lines to your `/etc/hosts` file:

AirLab servers

<magrathea_ip> magrathea magra

<elysium_ip> elysium ely

<westworld_ip> westworld ww

<multiverse_ip> multiverse multi

Now you can also use "nicknames" for the server (eg, `ww` instead of `westworld`)

SSH certificate

- Each time you connect to a server you need to put your credentials
- Tedious, error prone, need to remember a password

-> **use ssh certificate**

<https://linuxhint.com/generate-ssh-key-ubuntu/>

Do it for all the server you can connect to

SSH config

- Typing `ssh surname@server` each time is tedious as the username is always the same!
- We can configure ssh to do it for us!
- Add to your `.ssh/config` the following lines:

```
Host ely
  HostName ely
  User YOUR_USERNAME_HERE
  Port 22
```

```
Host magra
  HostName magra
  User YOUR_USERNAME_HERE
  Port 22
```

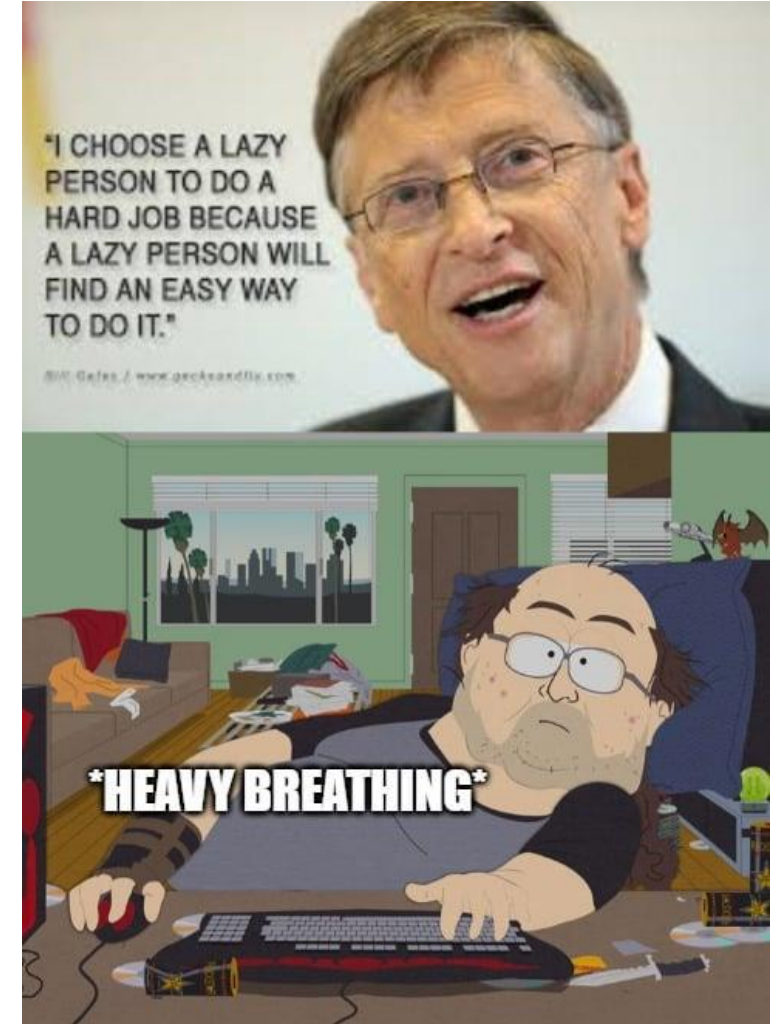
```
Host ww
  HostName ww
  User YOUR_USERNAME_HERE
  Port 22
```

ssh config

Finally now you can connect to any server with just:

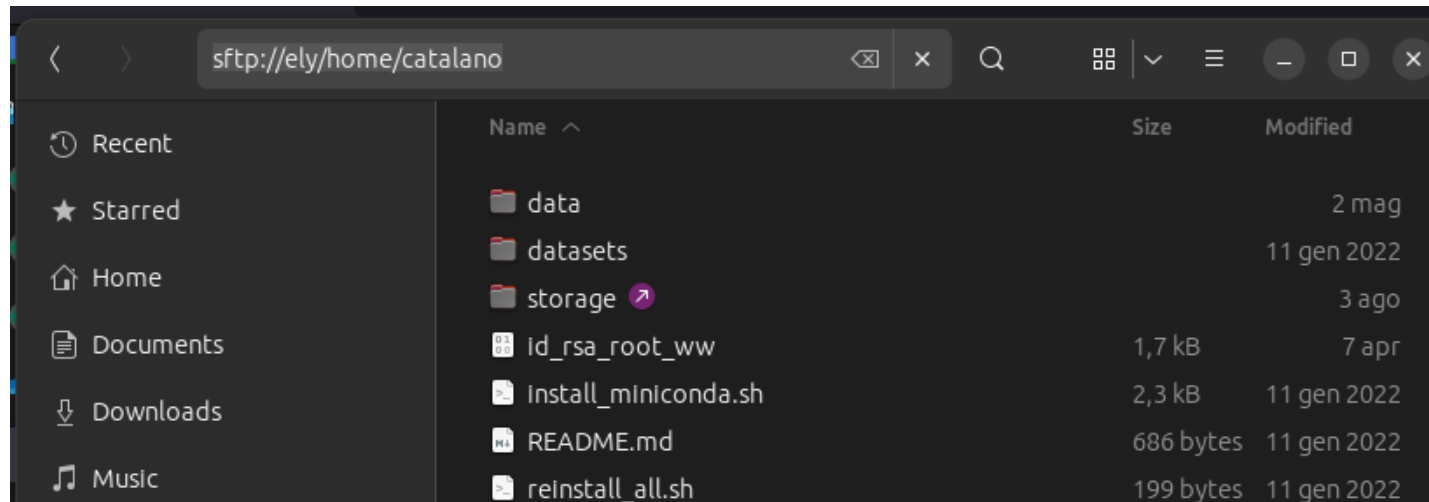
```
`ssh server_nickname`
```

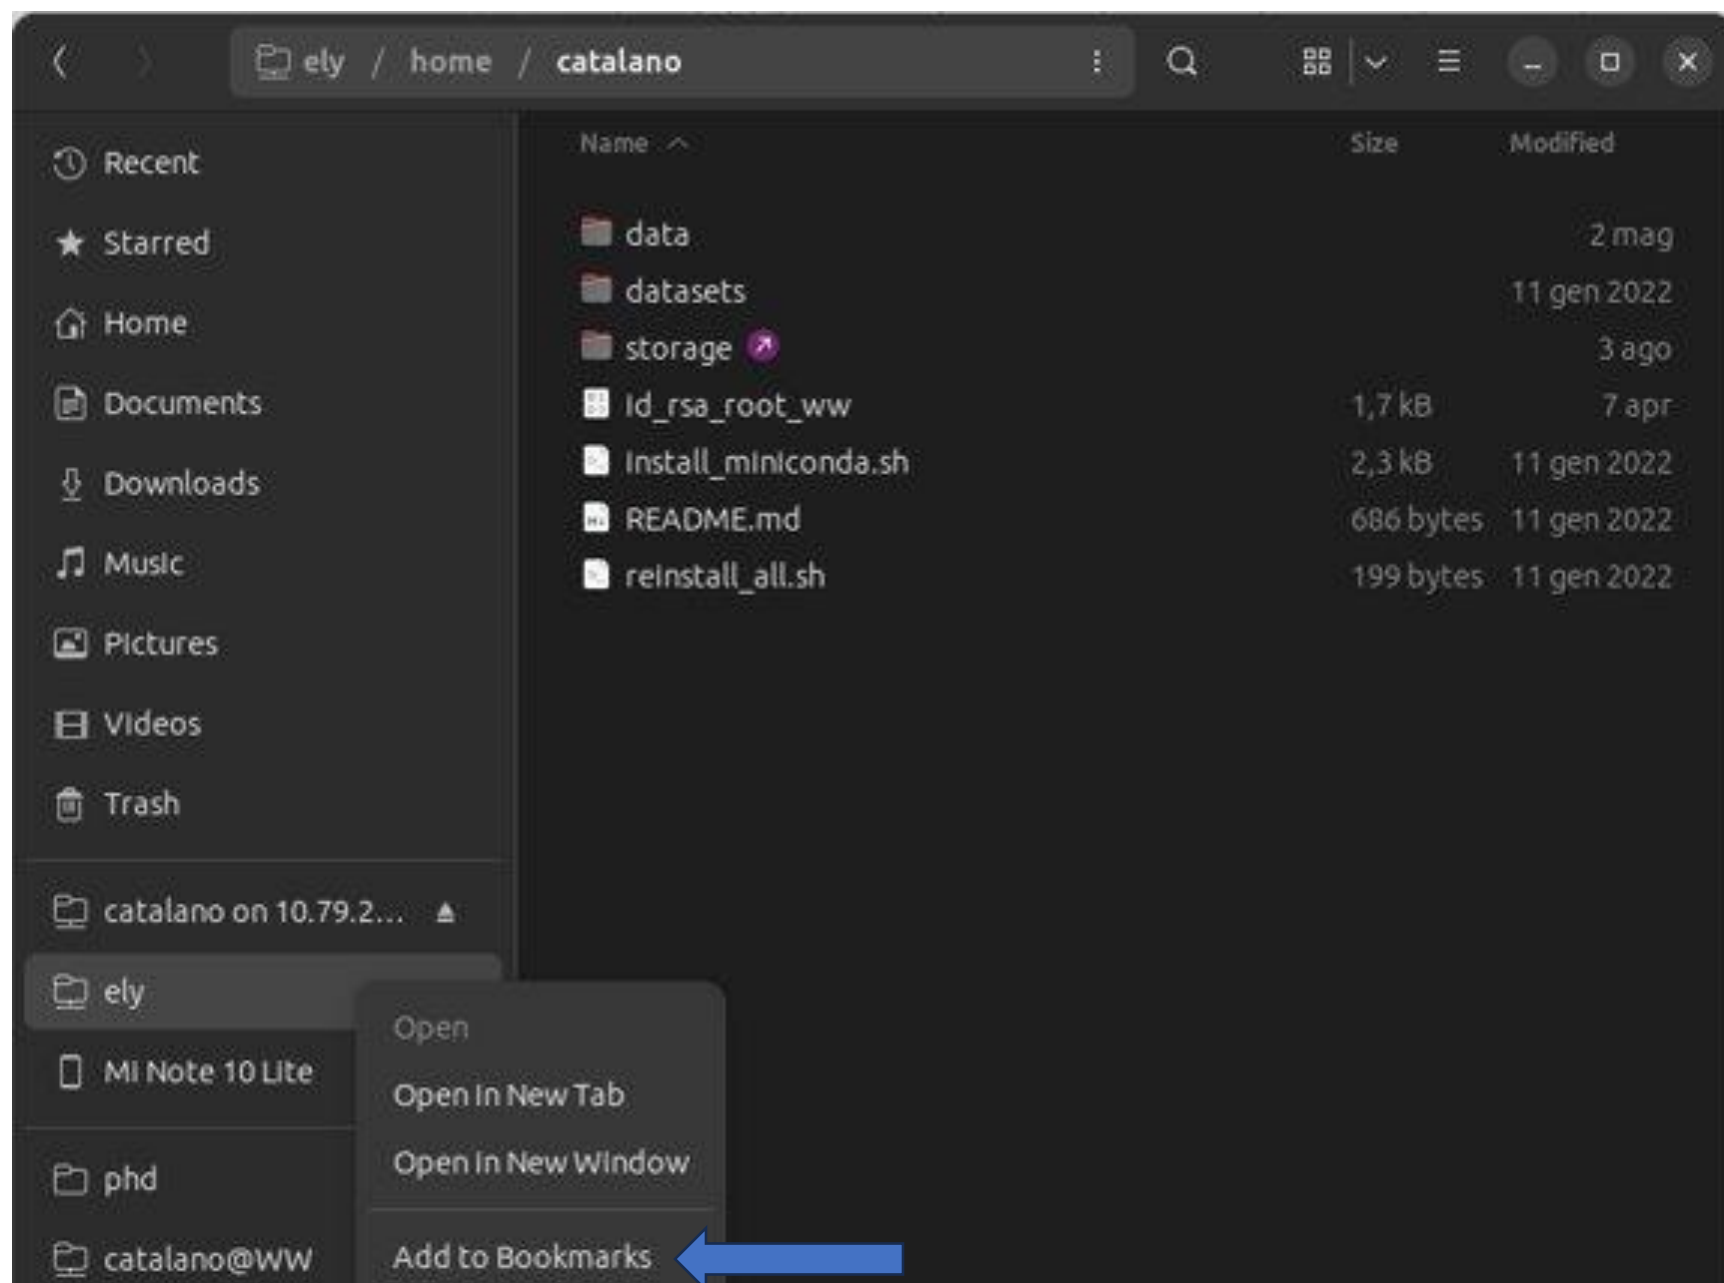
Bonus: you should now also have tab-autocomplete for server_nickname




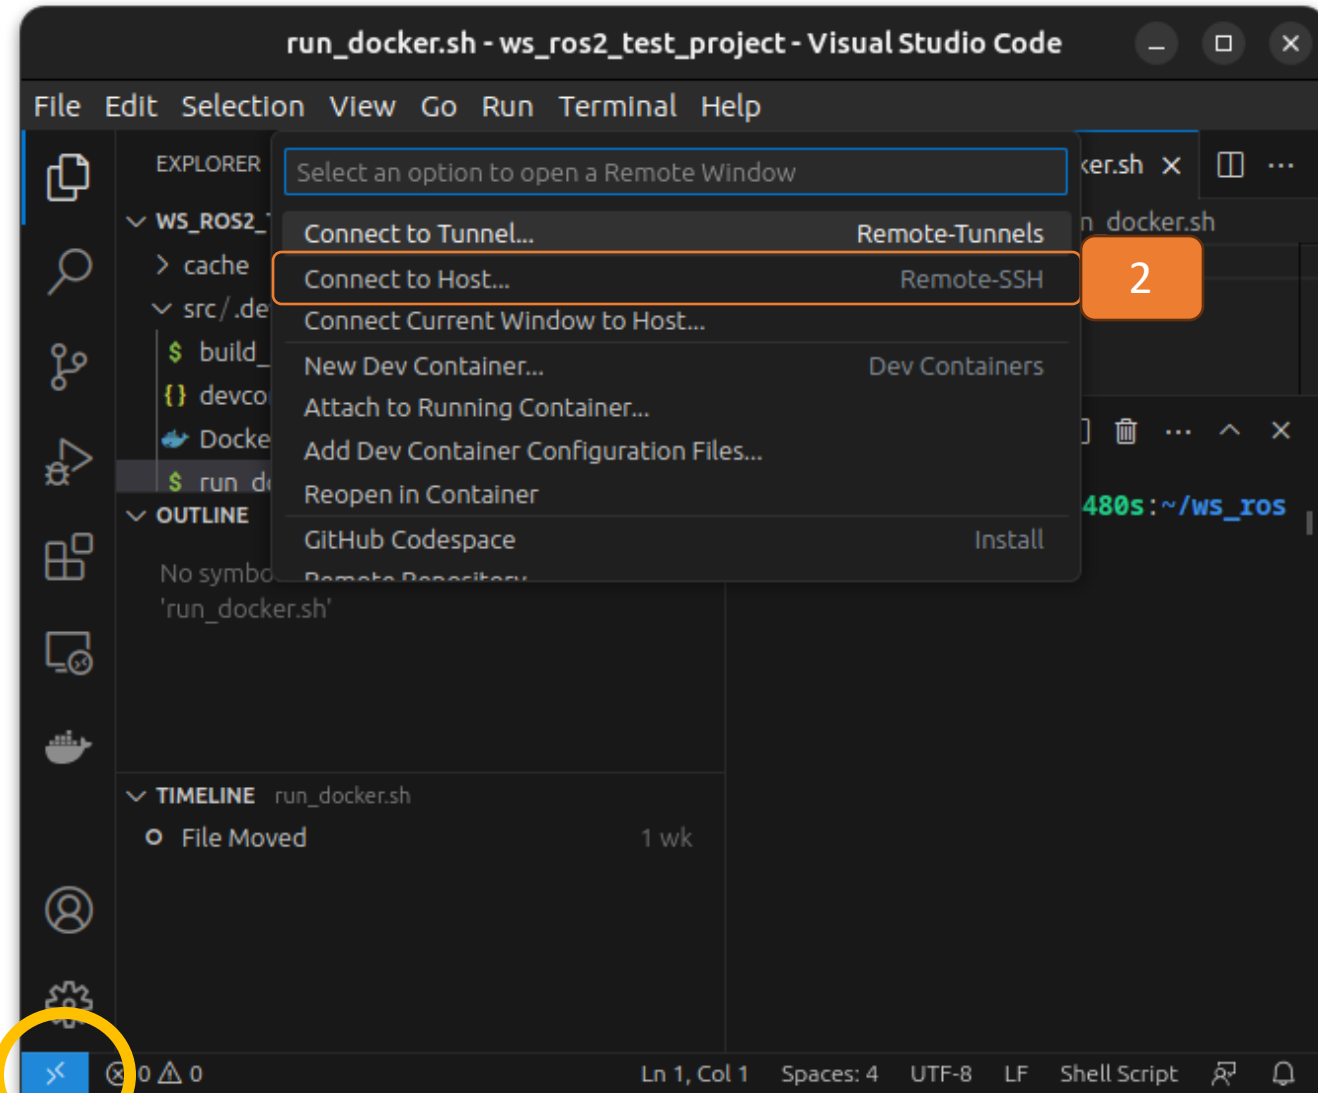
Move files from/to servers

- You can use FTP, scp, and rsync
- But you can also map your home on the server in nautilus on your local machine
- Use `sftp://server_name/home/username` as address
- Save this location to bookmarks





Configure VS Code Remote SSH



Have a github repo for your project!

- Sync your work on github repo
- Following these instruction on the sever you are developing will let you use git hub

<https://docs.github.com/en/authentication/connecting-to-github-with-ssh/generating-a-new-ssh-key-and-adding-it-to-the-ssh-agent>

<https://docs.github.com/en/authentication/connecting-to-github-with-ssh/adding-a-new-ssh-key-to-your-github-account>



tmux: keeping ssh session alive when disconnect from server

- Running a time intensive script and want to turn off your pc?
- Want to login on the same shell from different location?

use tmux!

- Cheat sheet here: <https://tmuxcheatsheet.com/>

Check how to create, attach, detach and delete a tmux session

tmux: things to know

- Scroll a tmux pane:

“Ctrl + b” followed by “[” (then you can use the arrows)

- tmux copy and paste:

Enter ‘copy mode’ by pressing “CTRL + b”, [.

Use the arrow keys to go to the position from where you want to start copying. Press “CTRL + SPACE” to start copying.

Use arrow keys to go to the end of text you want to copy. Press “ALT + w” or “CTRL+ w” to copy into the tmux buffer.

Press “CTRL +b”,] to paste in a possibly different Tmux pane/window.

tmux alternatives??

Do you know tmux alternatives?

Please let me know!!



Check GPU status

- Use command `nvidia-smi`

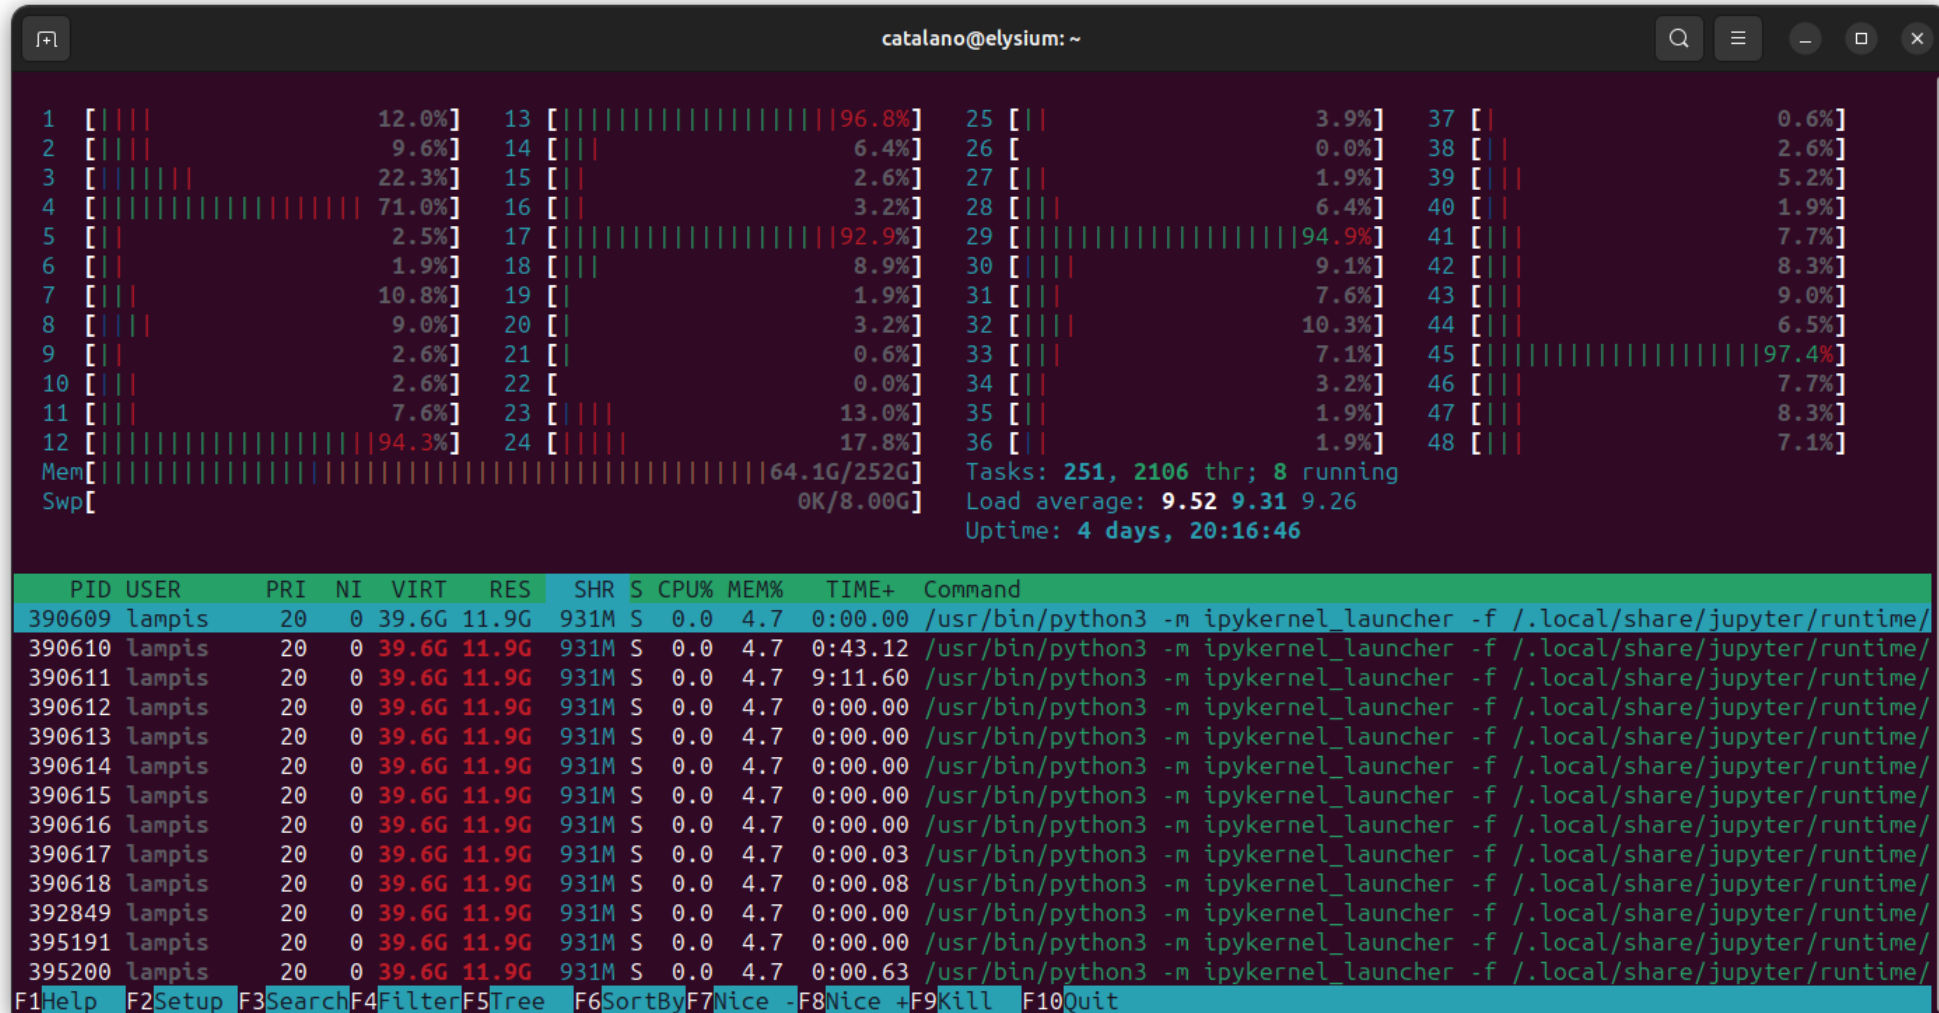
NVIDIA-SMI 520.56.06				Driver Version: 520.56.06				CUDA Version: 11.8			
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC					
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.					
						MIG M.					
0	Quadro RTX 6000	On	00000000:04:00.0	Off		Off					
74%	56C	P2	197W / 180W	9777MiB / 24576MiB	40%	Default					
						N/A					
1	Quadro RTX 6000	On	00000000:05:00.0	On		Off					
87%	60C	P2	99W / 180W	17301MiB / 24576MiB	0%	Default					
						N/A					
2	NVIDIA TITAN X ...	On	00000000:83:00.0	Off		N/A					
91%	66C	P2	179W / 180W	10051MiB / 12288MiB	100%	Default					
						N/A					
3	NVIDIA GeForce ...	On	00000000:84:00.0	Off		N/A					
32%	37C	P8	9W / 180W	13MiB / 11264MiB	0%	Default					
						N/A					
4	Quadro RTX 6000	On	00000000:87:00.0	Off		Off					
78%	57C	P2	151W / 180W	6655MiB / 24576MiB	74%	Default					
						N/A					
5	Quadro RTX 6000	On	00000000:88:00.0	Off		Off					
36%	40C	P8	14W / 180W	17301MiB / 24576MiB	0%	Default					
						N/A					
Processes:											
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage					
	ID	ID									
0	N/A	N/A	2558	G	/usr/lib/xorg/Xorg	8MiB					
0	N/A	N/A	14390	C	/usr/bin/python3	3122MiB					
0	N/A	N/A	2533867	C	python	6642MiB					
1	N/A	N/A	2559	G	/usr/lib/xorg/Xorg	8MiB					
1	N/A	N/A	388610	C	/usr/bin/python3	17288MiB					
2	N/A	N/A	2560	G	/usr/lib/xorg/Xorg	8MiB					
2	N/A	N/A	3829190	C	python	4314MiB					
2	N/A	N/A	3829728	C	python	4080MiB					
2	N/A	N/A	3840660	C	python	1644MiB					
3	N/A	N/A	2561	G	/usr/lib/xorg/Xorg	8MiB					
4	N/A	N/A	2562	G	/usr/lib/xorg/Xorg	8MiB					
4	N/A	N/A	2523223	C	python	6642MiB					

Get information about a process

- Use ``ps aux | grep PROCESS_PID``
- Useful to get the process' user

"task manager" for the command line

- Use `htop`




Developing pipeline - high level picture



Docker Image

The "virtual environment"
- what library to make
available to your script
- what dependences to install
- ..



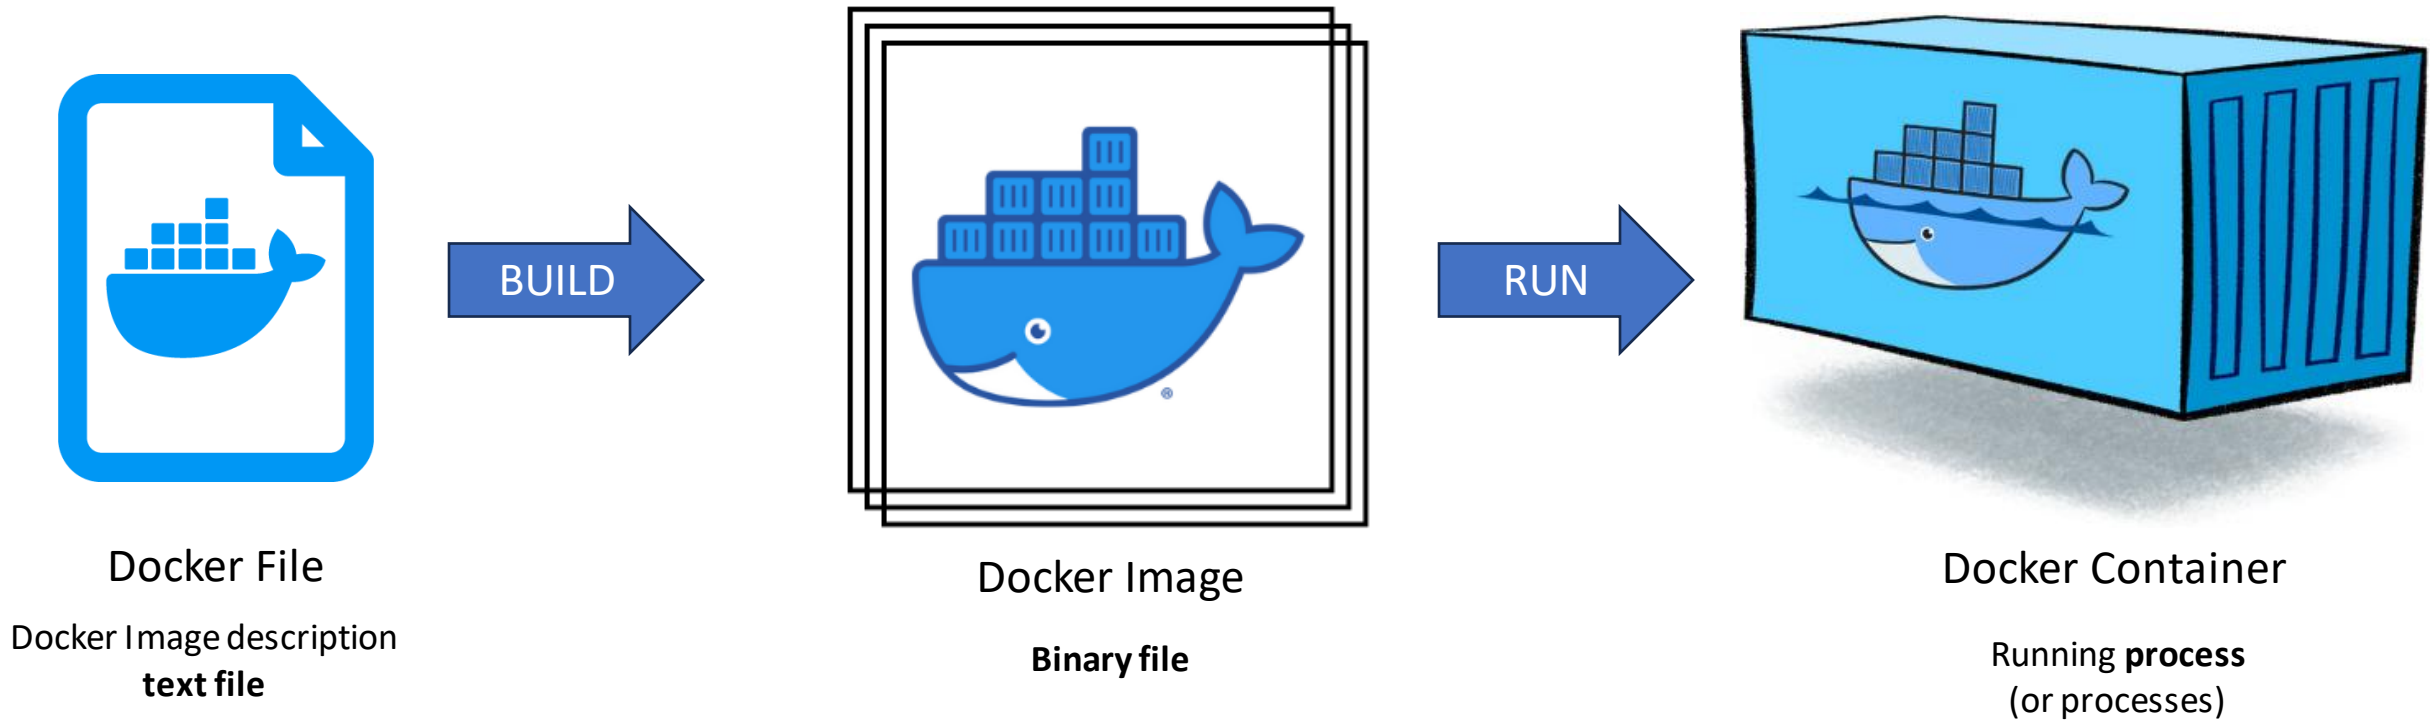
Docker Container

The image actually running. **It is immutable**

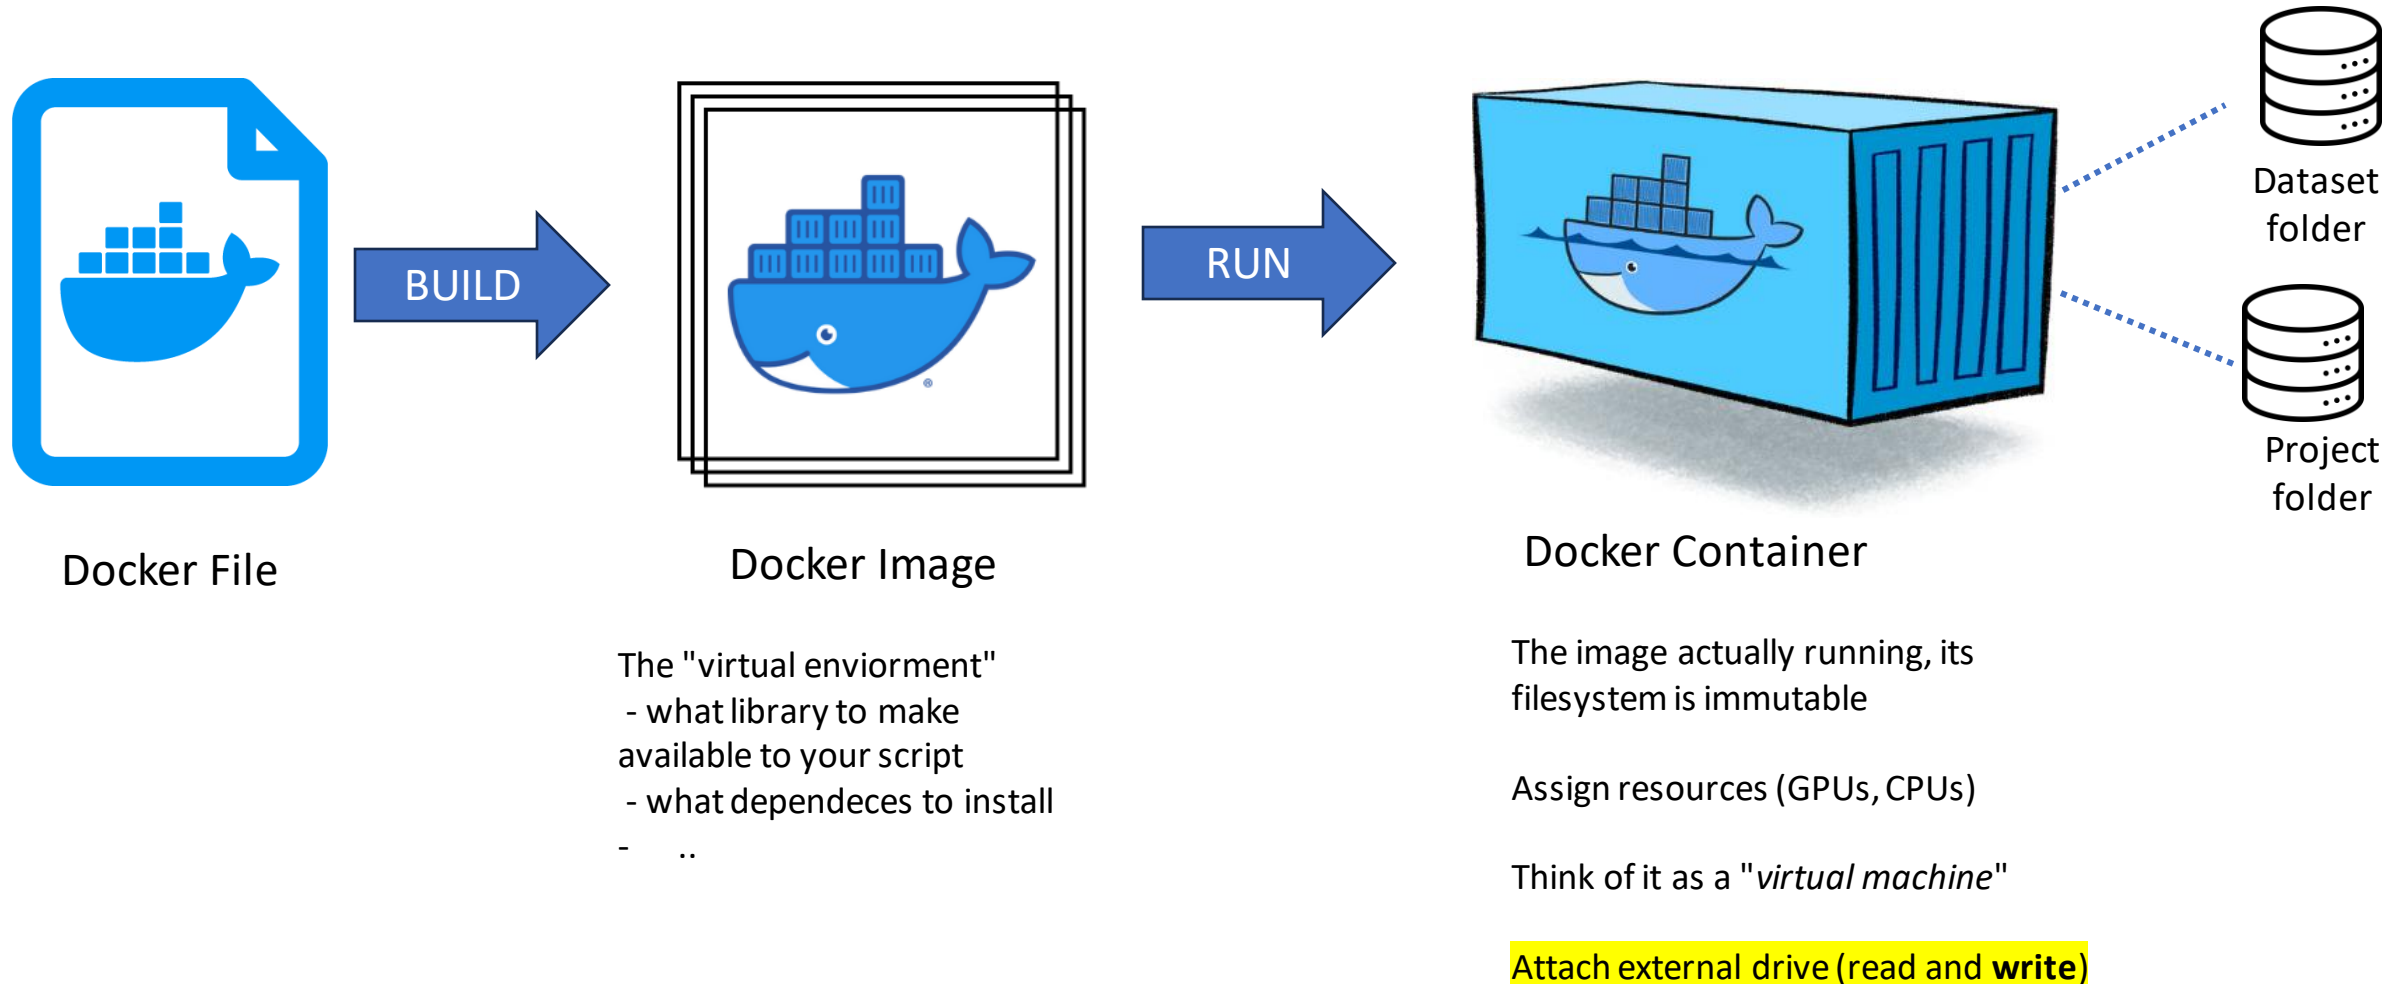
Assign resources (GPUs, CPUs)

Think of it as a "*virtual machine*"

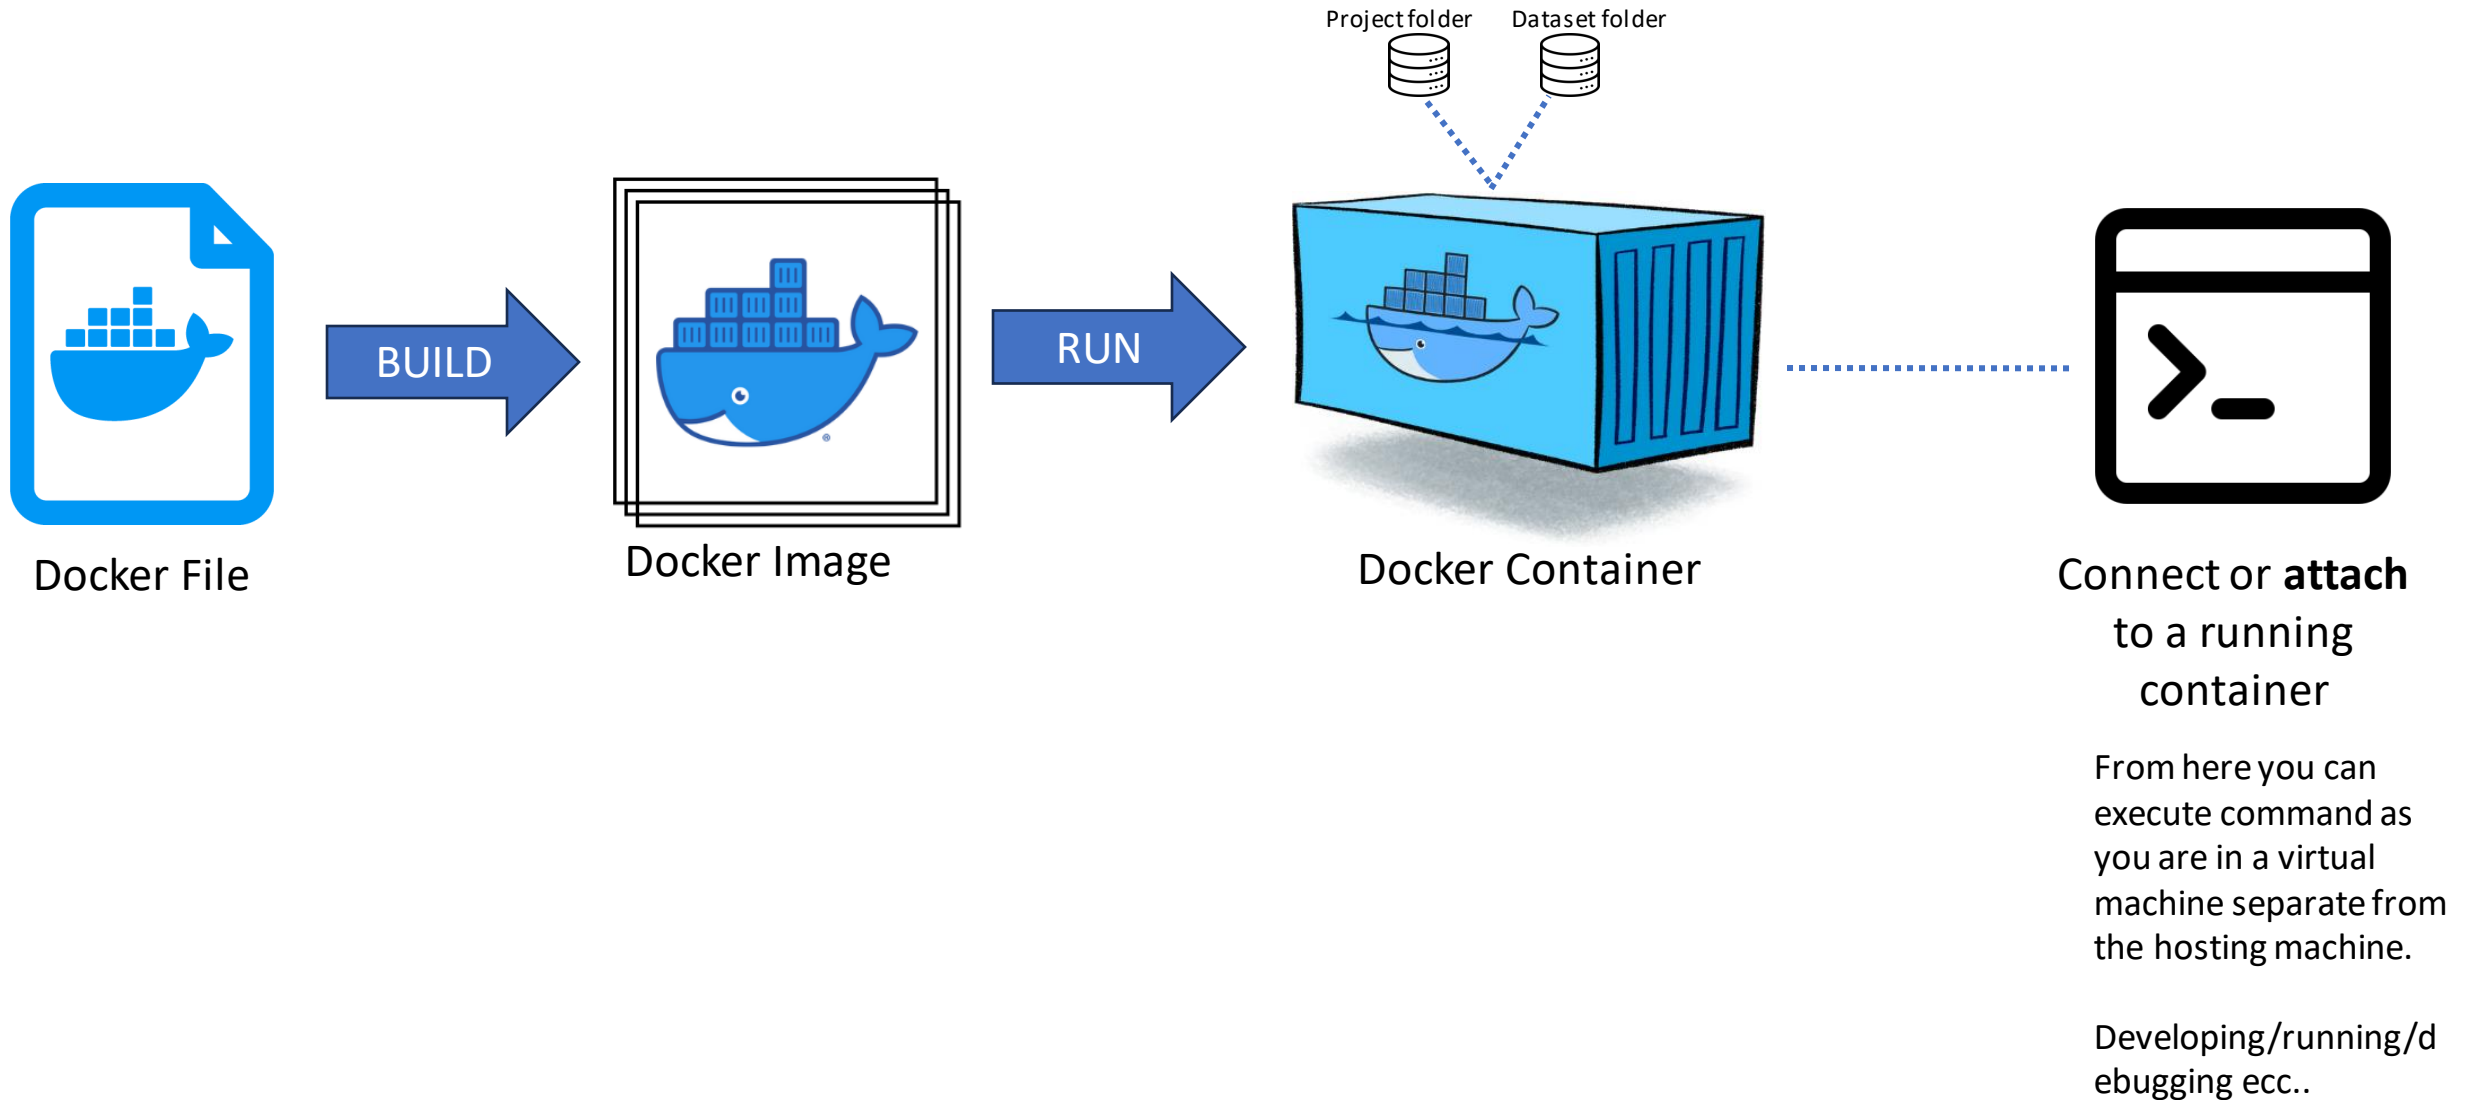
Developing pipeline - high level picture



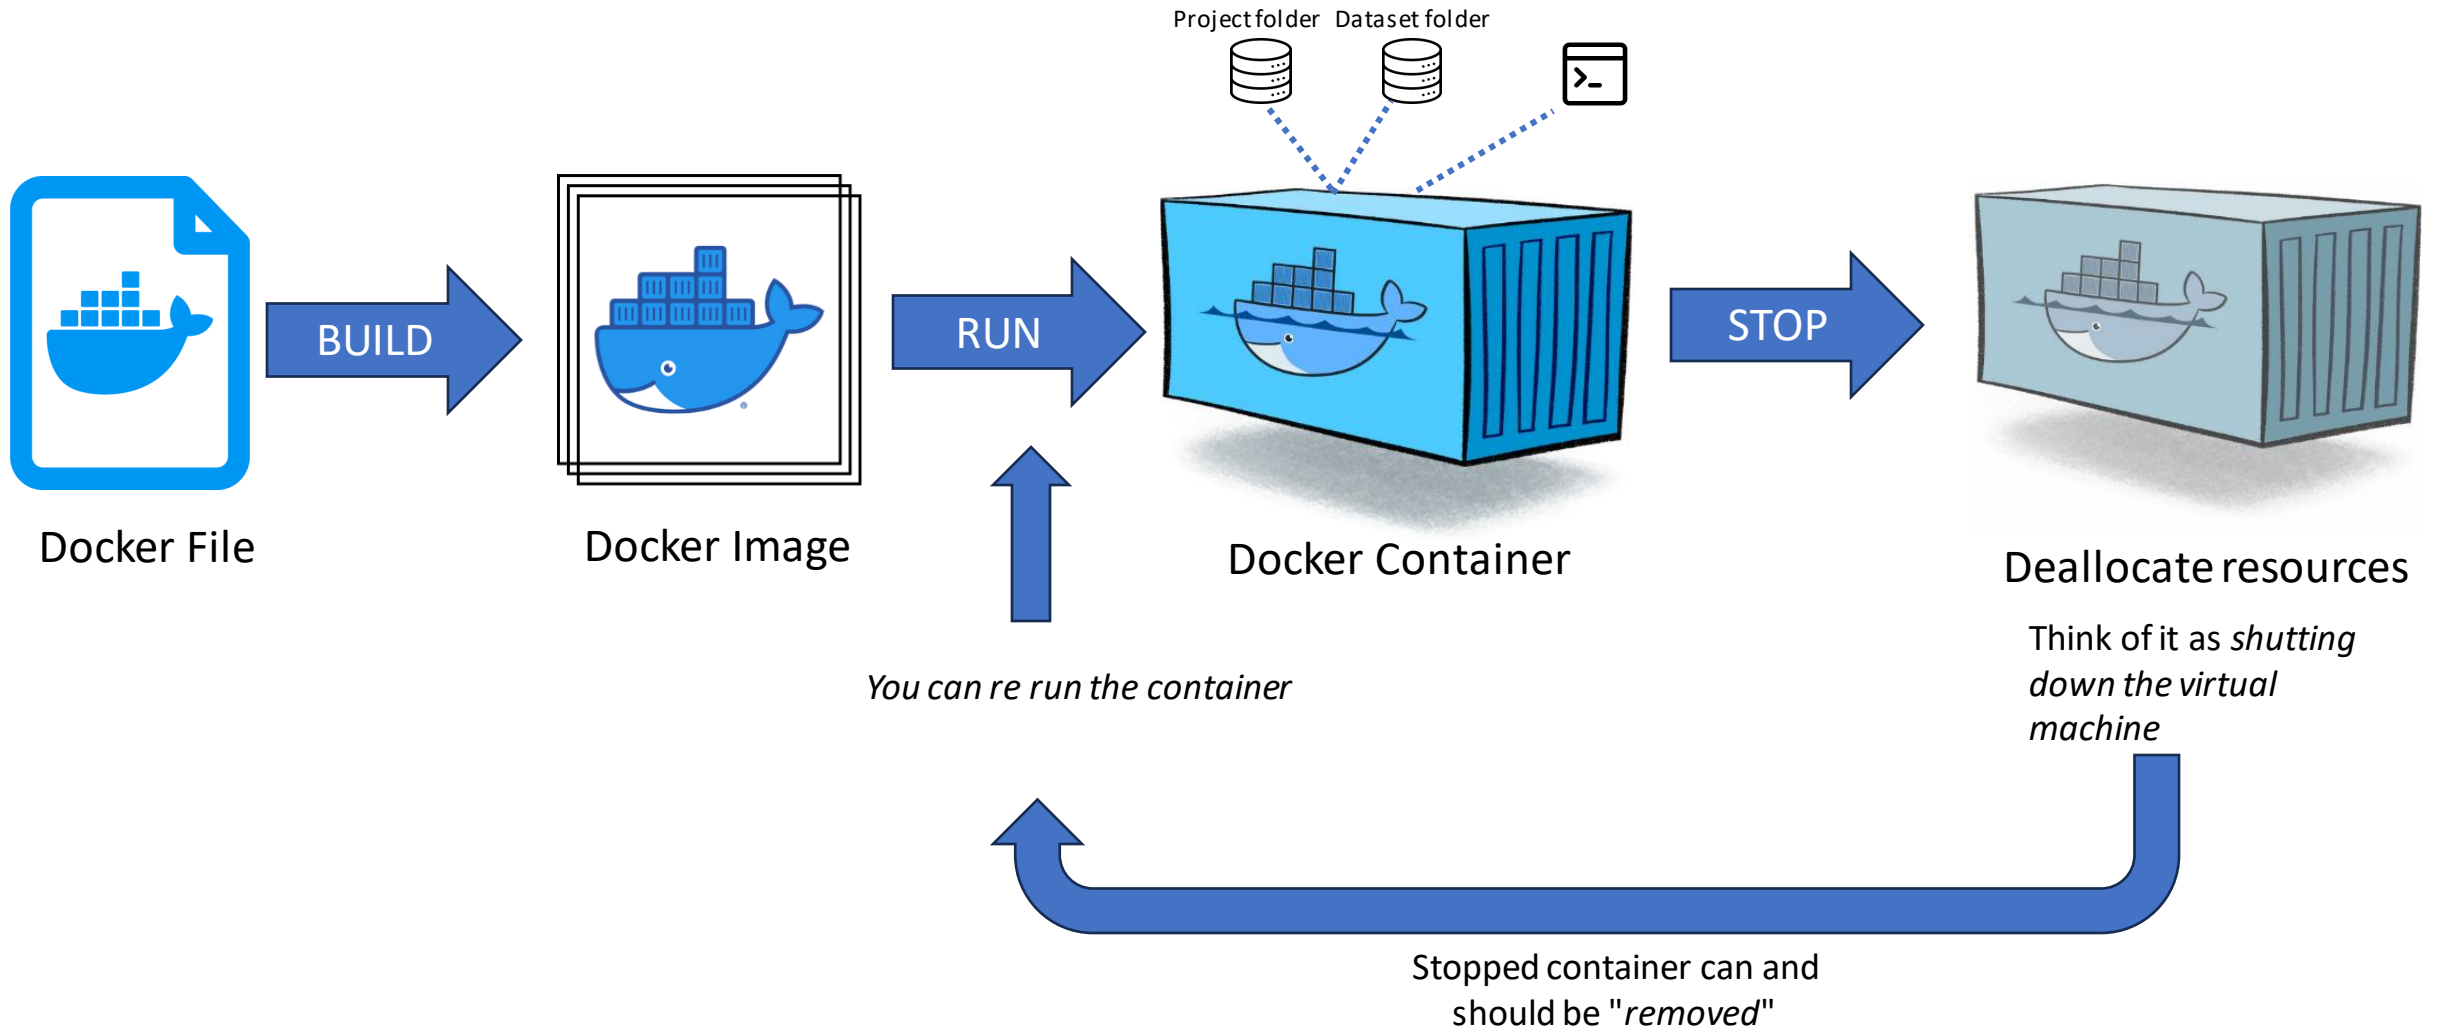
Developing pipeline - high level picture



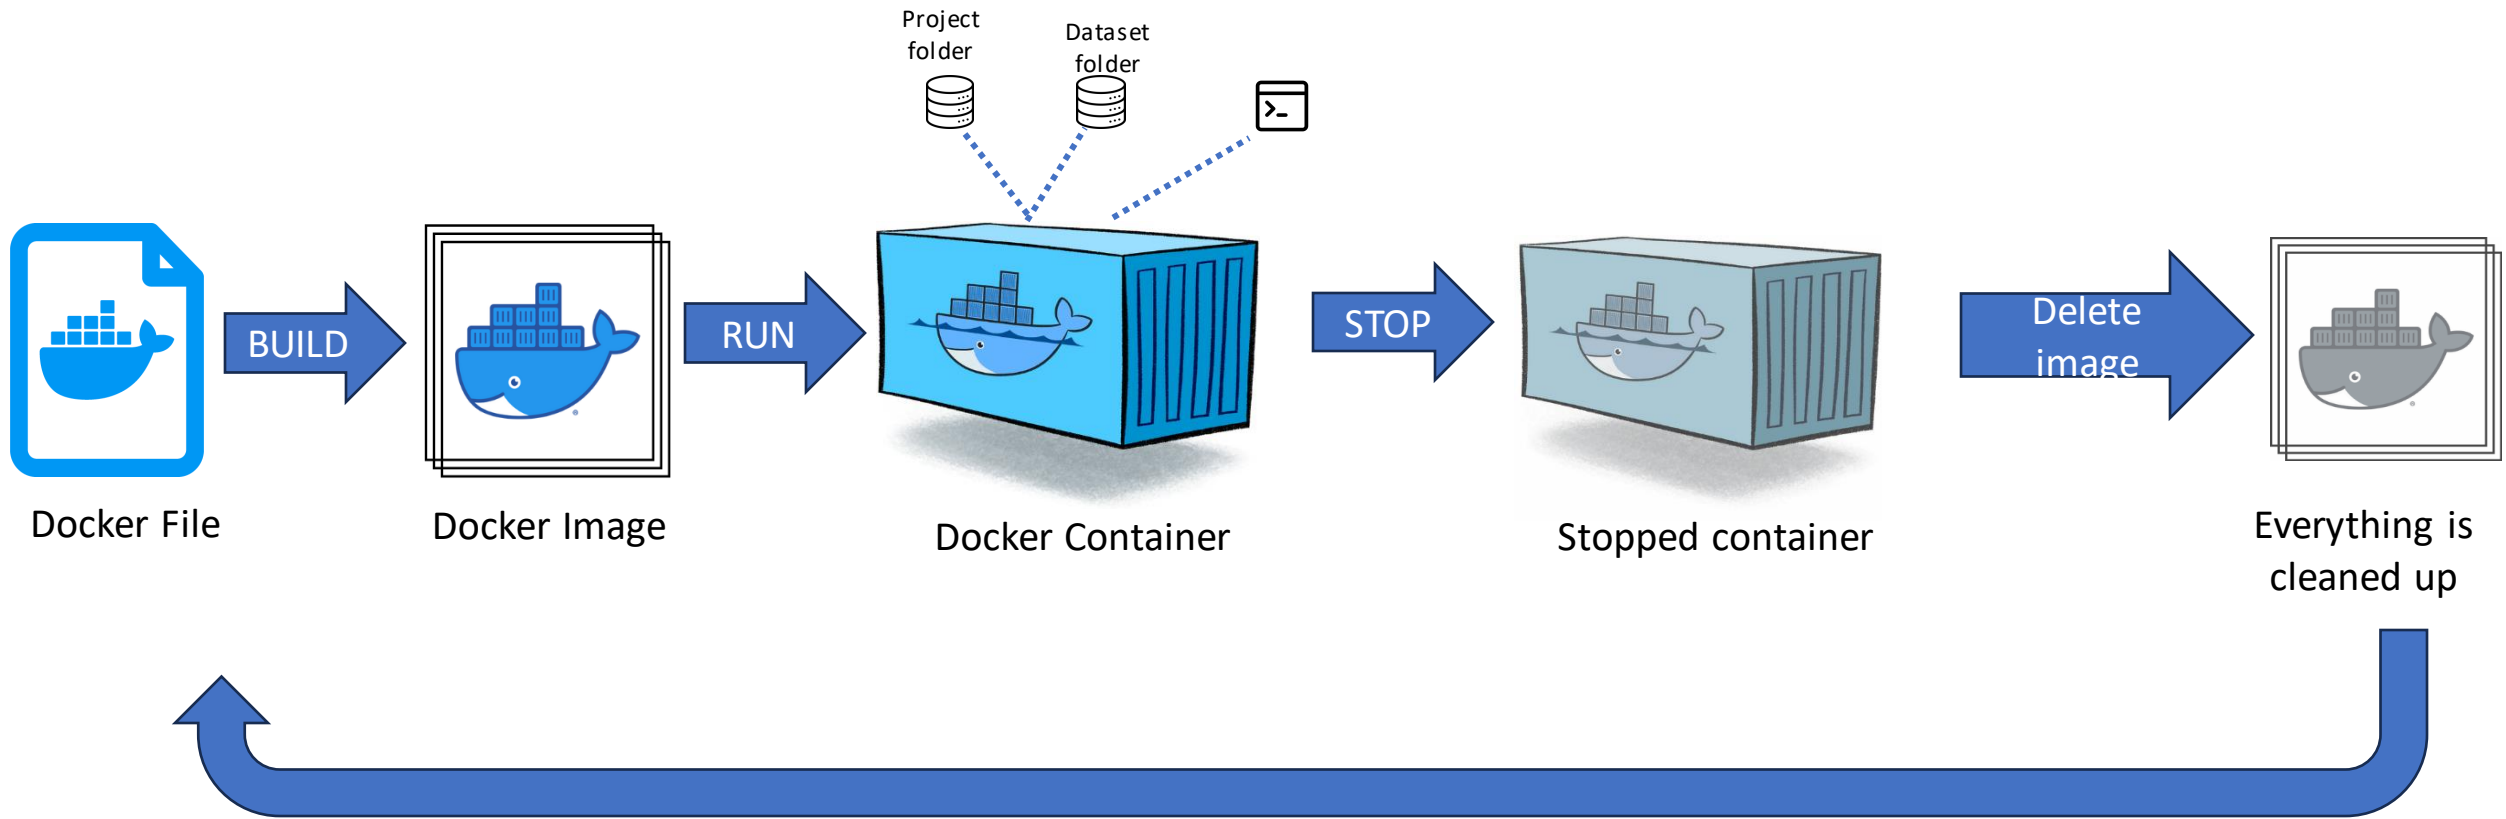
Developing pipeline - hig level picture



Developing pipeline - high level picture



Developing pipeline - high level picture



Docker useful commands

Build a new docker image:

`docker build --rm -t your_surname/image_name`

Visualize docker images:

`docker images`

`docker images | grep your_surname`

Run container

`docker run`

with some sugar

`run-docker`

Visualize running containers:

`docker ps`

`docker ps | grep your_surname`

Attach to a running container:

`docker exec -it CONTAINER_ID bash`

Stop container:

`docker stop CONTAINER_ID`

Remove container:

`docker rm CONTAINER_ID`

Remove a docker image:

`docker rmi your_surname/image_name`

Example of docker run

- You need to specify your UID and GID
 - Use command ``id surname`` to get it
- You need to specify which image to use
- You need to specify the resulting docker container name

Example of docker run

- Set UID and GID as 1085
- Set container name surname_docker_example
- Image to use surname/example_image:v0

```
docker run -e HOST_UID=1085 -e HOST_GID=1085 --name surname_docker_example surname/example_image:v0
```

Example of docker run

- Use GPU 7
- Use CPU from 104 to 111

```
docker run --gpus device=7 --cpuset-cpus 104-111 -e HOST_UID=1085 -e HOST_GID=1085 -u 1085:20002 --name surname_docker_example surname/example_image:v0
```

Example of docker run

- Set the working directory as /home/surname/exp

```
docker run --gpus device=7 --cpuset-cpus 104-111 -w /home/catalano/exp -e HOST_UID=1085 -e HOST_GID=1085 -u 1085:20002 --name surname_docker_example surname/example_image:v0
```

Example of docker run

- Mount the dataset folder
- Mount the project folder

```
docker run --gpus device=7 --cpuset-cpus 104-111 --mount  
type=bind,source=/home/surname/storage/project,target=/home/surname/exp --mount  
type=bind,source=/multiverse/datasets/surname/,target=/home/surname/exp/private_datasets -w /home/catalano/exp -e  
HOST_UID=1085 -e HOST_GID=1085 -u 1085:20002 --name surname_docker_example surname/example_image:v0
```


Example of docker run

- Use parameter `-d` to let the container run in the background
- Use parameter `--rm` to automatically remove the container when it exits
- Use parameter `-it` to have a console

```
docker run -d --rm -it --gpus device=7 --cpuset-cpus 104-111 --mount  
type=bind,source=/home/surname/storage/project,target=/home/surname/exp --mount  
type=bind,source=/multiverse/datasets/surname/,target=/home/surname/exp/private_datasets -w /home/catalano/exp -e  
HOST_UID=1085 -e HOST_GID=1085 -u 1085:20002 --name surname_docker_example surname/example_image:v0
```

Other parameters

- You might also want to :
 - set a memory limit
 - Forward some TCP port outside the container
 - ..
- Google/ChatGPT it

The lazy way

- Write a docker file starting from an existing one
- Have a build script
- Run the container via a bash script
- Attach your IDE to the running container
 - In VS Code you can:
 - Forward ports from a running container
 - Run Jupiter notebooks in the attached container w/o using the browser
- When you finish, stop the container, remove it, remove the image.
- Check your IDE plugins, there might be something useful for handling docker

Toy example

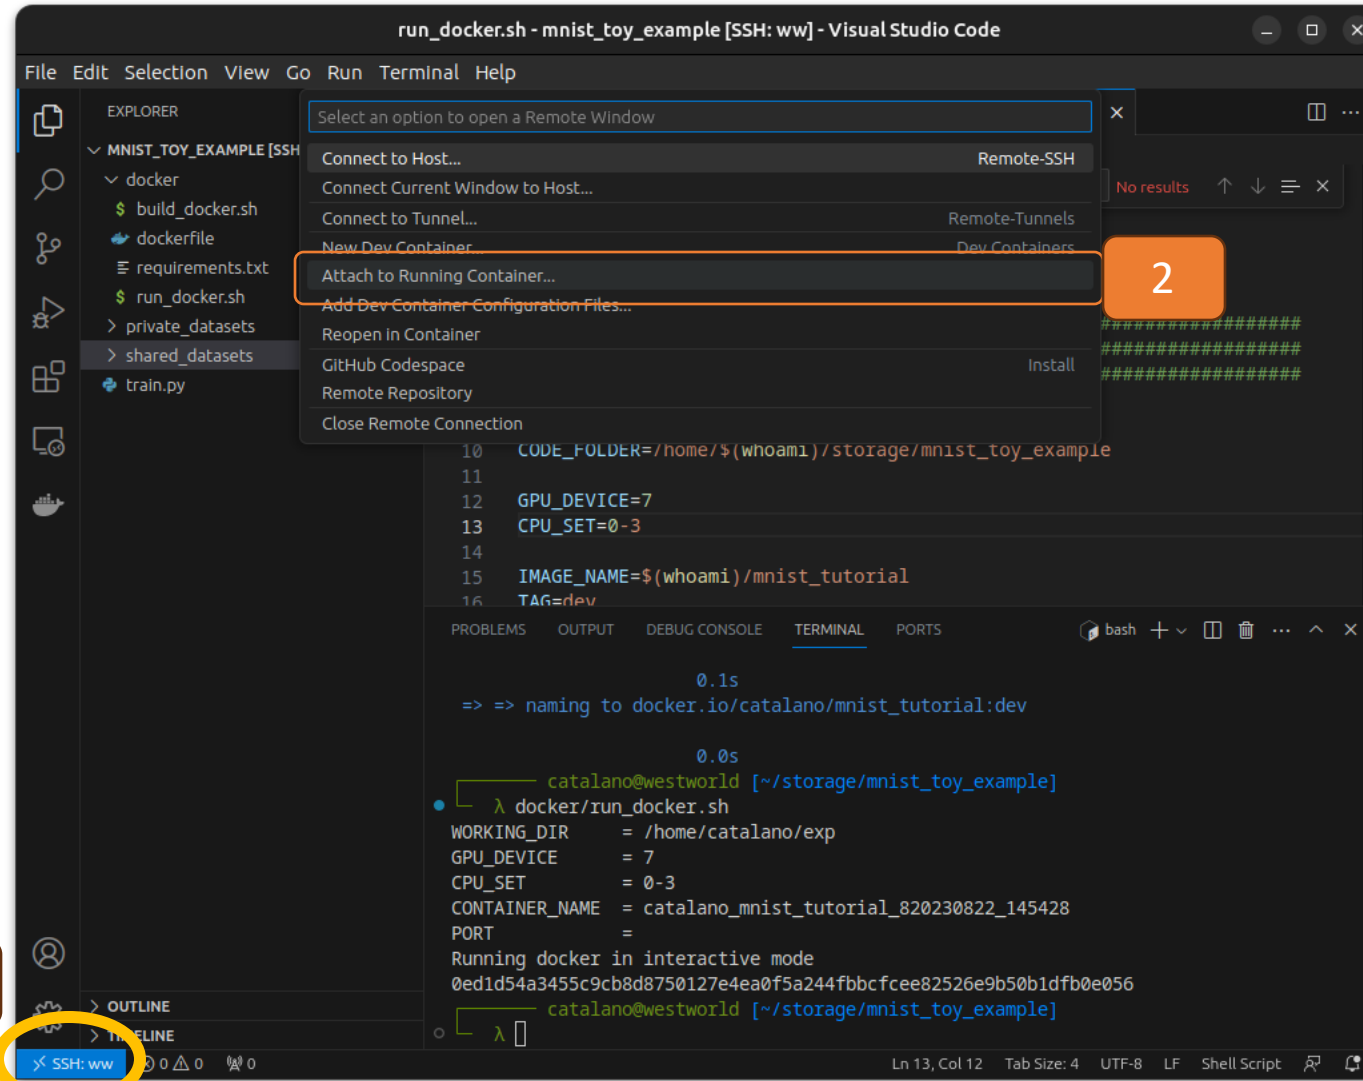
- MNIST classification example:

Repo on github

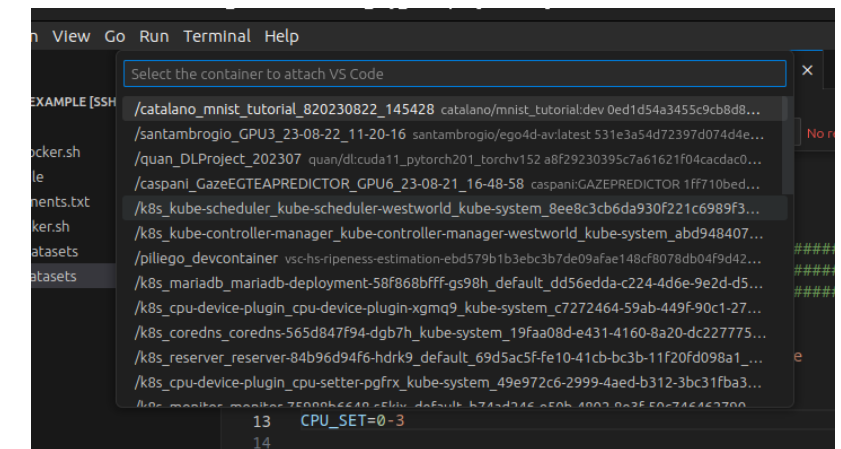
https://github.com/AIRLab-POLIMI/MNIST_Toy_Example

- We create the environment with the docker file, and install the python libs in requirements.txt
- Build the image with docker/build_docker.sh (it takes some time)
- Book the resources needed
- Run the container with
- Attach VS Code to the running container

Attach VS Code to an already running container



3 Select your container from the list



- Your project will be in /home/surname/exp
- CTRL + j to open a terminal
- **Now you can play**



Tensorboard

- Store and compare logs from many experiments
- Can visualize many kind of data:
 - Scalars:
 - Loss value, accuracy ..
 - Images
 - Model architecture
 - Embeddings

Tensorboard

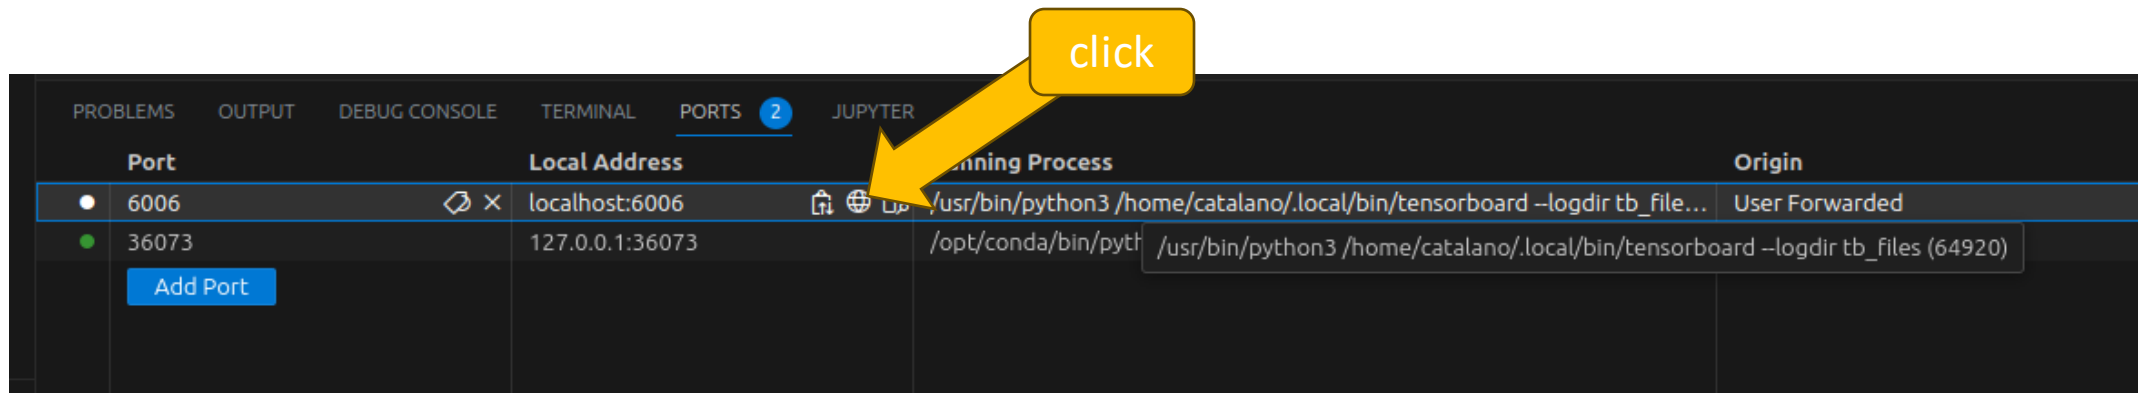
1. Create a folder for all your experiments
2. For each experiment create a "writer" with a meaningful name
3. Add values to the writer
4. To visualize the data run tensorbord with command:

```
tensorboard --logdir TB_DIRECTORY
```

By default it will listen on port 6006

Port forwarding

- When you run **tensorboard** VSCode should see that and auto forward its port outside.
- If not, or if you want to forward something else you can force the port forwarding
- ALT + j to open the bottom panel, go to the PORTS tab
- "Add Port"

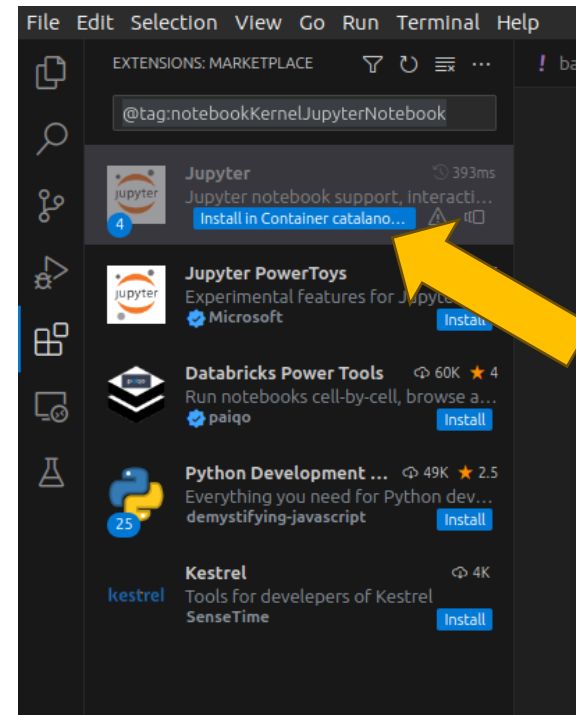
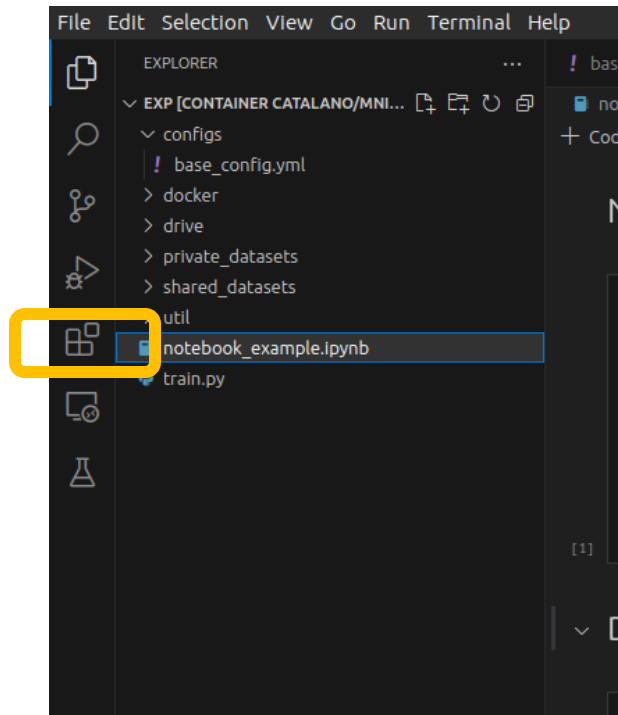


Tensorboard - Projector

- Let you project (PCA, UMAP,t-SNE) high dimensional object to a 2D or 3D space
- -> allow to have an idea on how "far apart" objects are in the embedding space
- For some scenarion can be extremly insigthfull (eg let you see outliers)
- Not well documented, but for each embedding you can attach some metadata (images, text, numbers, boolean) to do filters.
- Check example

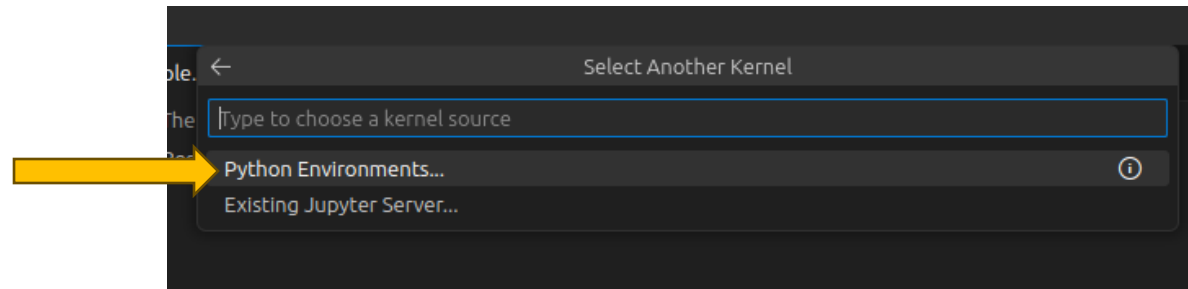
Jupyter notebook

- Install the extension @tag:notebookKernelJupyterNotebook



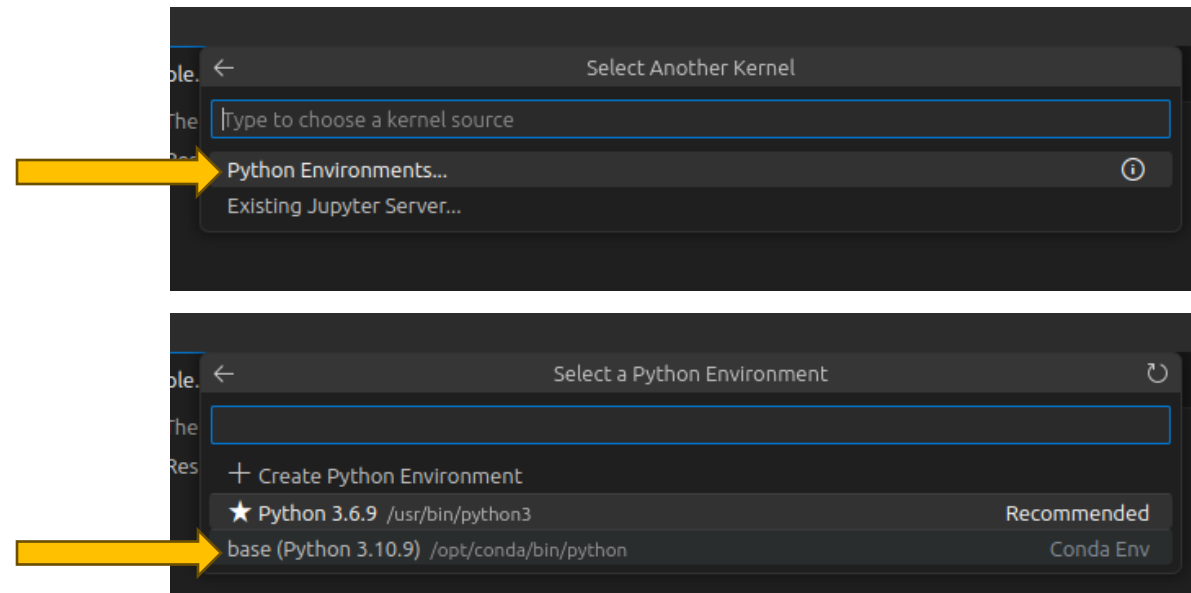
Jupyter notebook

- Install the extension @tag:notebookKernelJupyterNotebook
- Create/Open a ipynb file
- Select a kernel
- "python enviornment"



Jupyter notebook

- Install the extension @tag:notebookKernelJupyterNotebook
- Create/Open a ipynb file
- Select a kernel
- "python enviornment"
- "base"



Endorsed practice

- Use configuration files and keep the same source code to implement small variations in your experiments (use flags and values from a config file)
- Use tensorboard to log values
 - Losses
 - Accuracy
 - Ecc ..
- Use a log file instead of printing to the console
- Most importantly

Always code as if the guy who ends up maintaining
your code will be a violent psychopath who knows
where you live.

